

Introduction to Network Programming: Communication Endpoints

- Network Application Programming Interfaces (Network API)

- Application Programming Interface:

- **Application Programming Interface (API)** = is a collection of **library functions** through which **application programmers** (=you) can gain **access** to the **services** (provided by the system)

- Network APIs:

- The **network communication services** are accessed through **2 commonly used network APIs**:
 - **Berkeley Sockets API**
 - **(ATT's) Transport Layer Interface (TLI)** --- has been **superseded** by **XTI**

- Today:

- **Berkeley sockets** is now **by far the most popular Network API**

(We will learn **Berkeley sockets**....)

- Communication end points

- Communication endpoint:

- **Communication endpoint** = a **logical location** where **messages** are **transmitted/received**

- **2 kinds** of communication endpoints

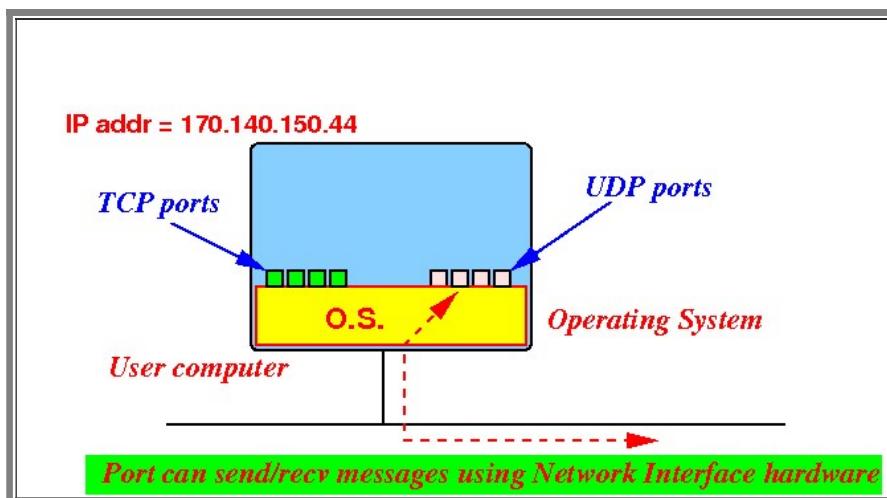
- **User program endpoint** = an **object** used in an **user program** that represents a **communication endpoint**
 - **(Operating) System endpoint** = an **object** used in the **Operating System** that represents a **communication endpoint**

- Communication end-points in the **Operating System**

- Ports:

- **Port** = a **communication end-point** inside the **Operating System**
 - **Ports** can make use of the **network interface card** of the **computer** to transmit **messages**

Graphically:



- Type of ports

- Types of **ports**:

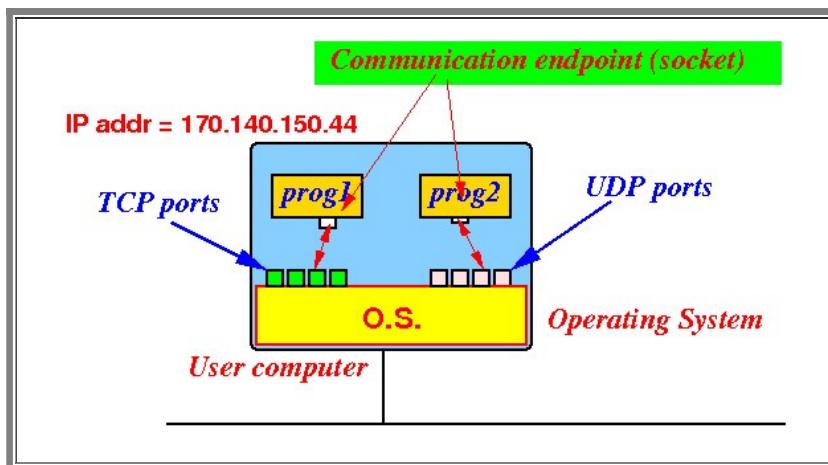
- The **Operating System** commonly implements **2 different types** of **ports** that provide **different type** of **communication services**
- **User Datagram Protocol (UDP) ports:**
 - **UDP ports** provides ***unreliable*** communication service
- **Transport Control Protocol (TCP) ports:**
 - **TCP ports** provides ***reliable*** communication service

- Communication end-points in a **computer program**

- **Socket:**

- **Socket** = a **communication endpoint** in a **network program**
 - The **network application program** can **send and receive messages** through **sockets**
- A **socket** must be **associated (bound)** to some **port** inside the **Operating System !!!**

- Network programming model in **computer programs**



Explanation:

- A **socket** can be **bound** to some **port** in the **Operating System**
- The **computer program** will use the **socket** to transmit message through the **port** to which the **socket** was **bound**

- Types of Sockets

- **Different types** of **sockets**:

- There are **different types** of **sockets**
- Each **type** of **socket** provides a specific **quality of service**

- There are **two types** of **sockets** in use today

- **UDP sockets**
 - A **UDP socket** is **bound** to an **UDP port**
 - **UDP sockets** will therefore provide ***unreliable*** communication service
 - The **UDP socket** uses a **packet-oriented protocol**: a **unit of transmission** is a **packet**

- **TCP sockets**

- A **TCP socket** is **bound** to an **TCP port**
- **TCP sockets** will therefore provide **reliable communication** service
- **TCP socket** uses a **byte-orientated protocol**: the **unit** of **transmission** is a **byte**

Introduction to Network Programming: socket variables

- Sockets inside a *computer program*: **socket variable**

- Recall:

- A **socket** is an **end-point of communication**
 - **Network application program** can **send and receive messages** through a **socket**

- **Socket variables:** representing a **socket** inside a **program**

- **Socket variable** = a **variable** used in a **computer program** that **represents/identifies** a **socket**
 - **Implementation in UNIX:**
 - A **socket variable** is just an **integer variable**
 - The **value** of the **socket (integer) variable** is used to **identify** an **entry** in the **communication ports** (= a **data structure**) **inside the (UNIX) Operating System**

(If you had taken **CS450**: a **socket** is like a **UNIX file descriptor**)

Don't worry if you had not taken CS450; you can understand this material without it.

- Operations on a socket

- Operations on a **socket**:

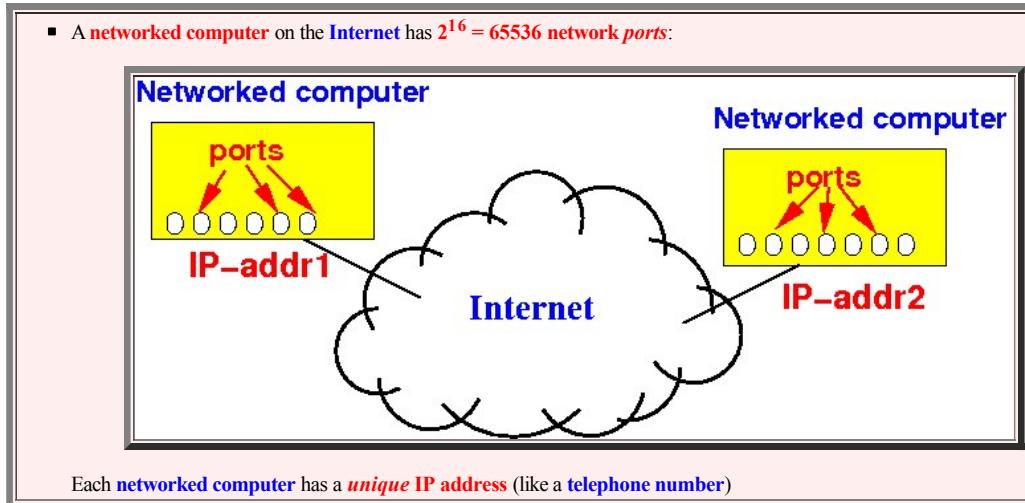
- You can perform a **write operation** on a **socket (variable)**
 - **writing** some **data** to a **socket (variable)** will **transmit** the **data** in a **message**
 - You can perform a **read operation** on a **socket (variable)**
 - **reading** a **socket (variable)** will **receive** a **message** (that was **transmitted** by some network program)

If there is **no message** in the **receive buffer**, the **read operation** will **block** until a **message** is **received**

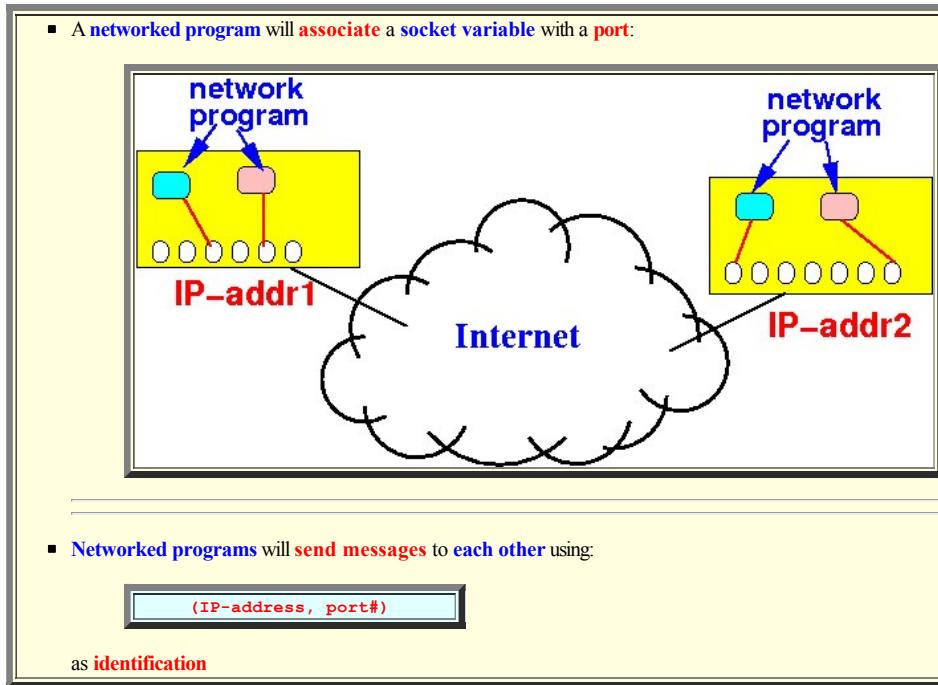
Overview network programming

- Overview network communication

- Infra-structure to provide network communication:



- Network programming:



Network Programming with UDP sockets

- C-language *Header files* for Network Programming

- Network programs written in C must include the following **header files** to define the **necessary data structures** used by the **Socket API**:

```
#include <stdio.h>           // Standard IO, all C programs need this
#include <stdlib.h>           // Standard Library, all C programs need this

#include <sys/types.h>         // System variable types
#include <sys/socket.h>         // Type definitions related to sockets
#include <netinet/in.h>         // Type definitions related to Internet protocol

#include <arpa/inet.h>          // Network utility functions
#include <netdb.h>              // Name services functions
```

Postscript - comments

- Advanced network programming: non-blocking sockets

- Default socket behavior for the **read operation**:

- The **read operation** on a **socket** by default is:

- **blocking**

- The program will **wait** until a **message** on the **socket** is **received**; then the **read operation** will **return** (with the **data** in the **newly received message**)

- Non-blocking socket:

- You can **change** a **socket** into **non-blocking** using the following **C-function call**:

- `fcntl(socket_variable, F_SETFL, fcntl(recvfd, F_GETFL) | O_NDELAY);`

- We will learn to use **blocking sockets** only....

- Network programming resources

- Socket programming guides:

- A very good introduction to network programming by L. Besaw: [click here](#).
 - A advanced socket API documentation: [click here](#).

Intro to UDP

- Properties of UDP communication
 - Properties of UDP communication:

- Block oriented:
 - A UDP message is transmitted as a *whole*
 - The receiver will receive a transmitted UDP message as a *whole*

I.e.:

- The receiver will receive what was sent
- The receiver will not receive a *portion* of the transmitted message

- Unreliable:
 - A message transmitted through UDP is "send" and "done" with no guarantee whatsoever

So:

- You can transmit a UDP message *even* when there is *no receiver* at the *receiving end* !!!

(It can be quite **disturbing** to novice users)

Creating a UDP socket

- Creating a socket in general

- Defining a **socket variable** in general:

```
int s;  
  
s = socket(domain, socket-type, protocol-used); // Creates a socket
```

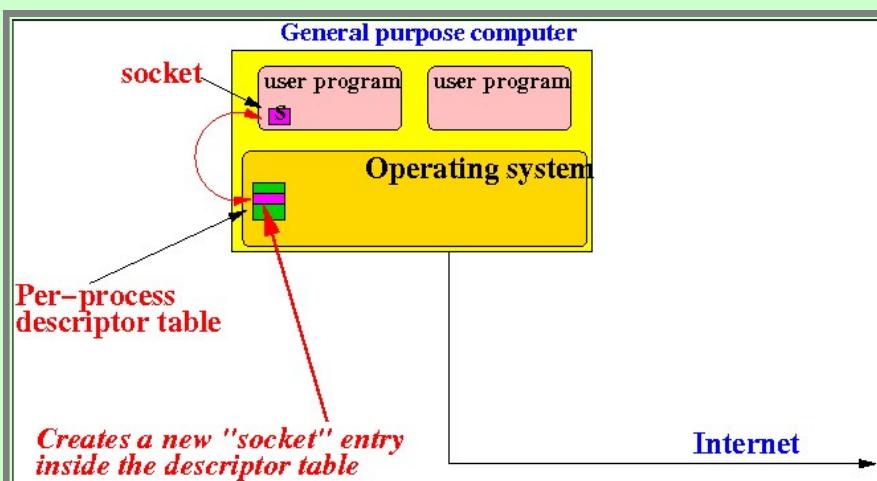
- Meaning of the **parameters**:

```
domain parameter:  
  
PF_INET = will use the IPv4 protocol (PF = protocol family)  
PF_INET6 = will use the IPv6 protocol  
  
socket-type parameter:  
  
SOCK_DGRAM = UDP socket  
SOCK_STREAM = TCP socket  
  
protocol parameter:  
  
used to select a network protocol that will/can provide the service  
0 = use the default network protocol  
  
SOCK_DGRAM's default protocol is: UDP  
SOCK_STREAM's default protocol is: TCP
```

- Result of the **socket system call**:

- The **socket()** system call will:
 - create an **entry** in the ("per-process") **descriptor table** inside the **Operating System**

Graphically:



- The **descriptor table** is used to:

- record the **state** of all **opened system resources** by the **process (running program)**

Example of system resources:

- Files
 - Pipes
 - Sockets
 - Etc., etc

- Creating a UDP socket

- System call to **create a UDP socket**:

```
int s;
```

```
s = socket(PF_INET, SOCK_DGRAM, 0); // Creates a IPv4 UDP socket
```

- **Associating a UDP socket with an UDP port (in the kernel)**

- **Fact:**

- After creating a **UDP socket**, we still need to *associate* the **socket** with an **UDP port** inside the **Operating System**

- **But first:**

- We need to **learn**:

- How to specify a **Internet network address** (that **identifies** a **UDP port**)

Identifying a host UDP endpoint (IP-Addr,Port#)

- The socket address structure for *Internet* programming in C

- Sockets:

- (Network) socket == a program endpoint for sending/receiving data
 - A socket is associated with a network port that is identified by a socket address
 - Socket address consists of:
 - IP address of the host
 - The port# of the kernel (UDP) port

- Data structure in C used to store Socket addresses of an (Internet) UDB port:

```
struct sockaddr_in
{
    sa_family_t      sin_family;    /* Socket's address family
                                         - Must contain the value AF_INET */

    in_port_t        sin_port;     /* Socket's Internet Port #
                                         - a 2 bytes port number */

    struct in_addr sin_addr;     /* Socket's Internet Address (IP-address)
                                         - a 4 bytes IP address */
};
```

Meaning of the variables:

- **sin_family**: (Internet socket family)
 - The **sin_family** must be set to the (constant) value **AF_INET** for the **IP protocol**
- **sin_port**: (the data type of this variable is **short** --- Internet socket port#)
 - If **sin_port > 0** then:
 - **sin_port** = the port# of some port in the **Operating System**
 - If **sin_port = 0** then:
 - the system will pick an **unassigned port** for you

Very important:

- **sin_port** must be stored in **network byte order** (more later) !!!!
- **sin_addr**: (the data type of this variable is a **struct** --- one of the field (**s_addr** contains the IP address))
 - **sin_addr** = the **4 byte IP address**
 - Note:
 - There is a special **wildcard IP address** named: **INADDR_ANY**
 - We will discuss this in the **multi-homed IP host** section later

Very important:

- **sin_addr** must also be stored in **network byte order** (more later) !!!!
- More on the **sin_addr** field:
 - The **sin_addr** variable is a **struct typed variable** used to store an **IP address**:

```
    uint32_t s_addr; /* 4 bytes IP address in network byte order */
} sin_addr;
```

- Preliminary example: how to initialize an the IPv4 address structure

```
(1) We first define a socket address variable:  
  
    struct sockaddr_in in_addr;  
  
(2) Then we initialize the socket address:  
  
    in_addr.sin_family = AF_INET; /* Protocol domain is: Internet */  
    in_addr.sin_addr.s_addr = ....; /* IP address in netw byte order */  
    in_addr.sin_port = ....; /* Port # in netw byte order */
```

(We will complete the example later after a discussion on byte ordering)

- Host byte ordering

- Host byte ordering:

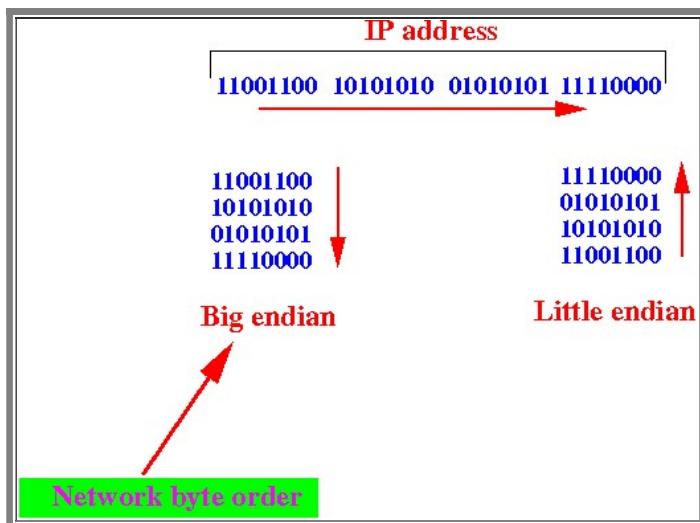
- Host byte ordering = the byte ordering used by the host computer

- Recall from CS255:

- A host computer can use one of 2 different ways to store serie of bytes in the memory of a computer:

- big endian and
 - little endian.

Example:



- Network byte ordering

- Network byte order:

- Network programming must always use the big endian order

- Specifically:

- These variables in the socket Internet address structure

- sin_port (Socket's Internet Port number)
 - sin_addr.s_addr (Socket's Internet Address)

- must use the network byte ordering !!!

- Utility functions to convert from host ordering to network ordering

- Converting from **host ordering** to **network ordering**:

■ **int htonl(int)**: (host to network long)

- **input = 4 bytes in host ordering**
- **output = 4 bytes in network ordering**

Example usage:

```
struct sockaddr_in    in_addr;

int   IP = 170*256*256*256 + 140*256*256 + 150*256 + 1; // 170.140.150.1

in_addr.sin_addr.s_addr = htonl(IP);
```

■ **short htons(short)**: (host to network long)

- **input = 2 bytes in host ordering**
- **output = 2 bytes in network ordering**

Example usage:

```
struct sockaddr_in    in_addr;

in_addr.sin_port = htons(8000);
```

- Example: setting up an **IPv4 socket address structure** with

- **IP address = 170.140.150.1**
- **Port # = 8000**

Program code:

```
/* =====
 Define a socket address variable
 ===== */
struct sockaddr_in    in_addr;

/* =====
 Initialize the socket address
 ===== */

in_addr.sin_family      = AF_INET; /* Protocol domain is: Internet */

int   IP = 170*256*256*256 + 140*256*256 + 150*256 + 1;
      // IP address is in host byte order !!!

in_addr.sin_addr.s_addr = htonl(IP); /* IP address in network byte order */
in_addr.sin_port        = htons(8000); /* Port # in network byte order */
```

- Utility functions to convert from **network ordering** to **host ordering**

- Fact:

- If we want to **print** the **IP address** or the **port #** stored in an **socket address variable**:

- We must **first convert** the **network byte order representation** to **host byte order** before **printing !!!**

- Converting from **network ordering** to **host ordering**:

■ **int ntohs(int)**: (network to host long)

- **input = 4 bytes in network ordering**
- **output = 4 bytes in host ordering**

Example usage:

```
struct sockaddr_in    in_addr;

int   IP;

IP = ntohs(in_addr.sin_addr.s_addr);

/* =====
 If you need to use (e.g., print) the value, you need to
 convert it to HOST order first
 ===== */
```

```
printf("%d", IP); // Prints the integer repr of the IP address  
// Computer program statements always use "host ordering" !!
```

- **short ntohs(short):**(host to network long)

- **input = 2 bytes in network ordering**
- **output = 2 bytes in host ordering**

Example usage:

```
struct sockaddr_in    in_addr;  
  
short                 portNum;  
  
portNum = ntohs(in_addr.sin_port);
```

Binding to a specific (kernel) UDP port#

- Binding a UDP socket

- Associating (= Binding) a *socket* variable to a given *network address*:

```
int bind(int s, (struct sockaddr *)addr, int length)

    s      = socket variable
    addr   = a socket IP address structure
    length = length of the socket IP address structure

Return value:

    0          = success
    Non-zero value = error code
```

Example:

```
int    IPAddress;           (a 32 bit IP address)
short  PortNumber;         (a 16 bit port number)

int    s;
struct sockaddr_in in_addr;

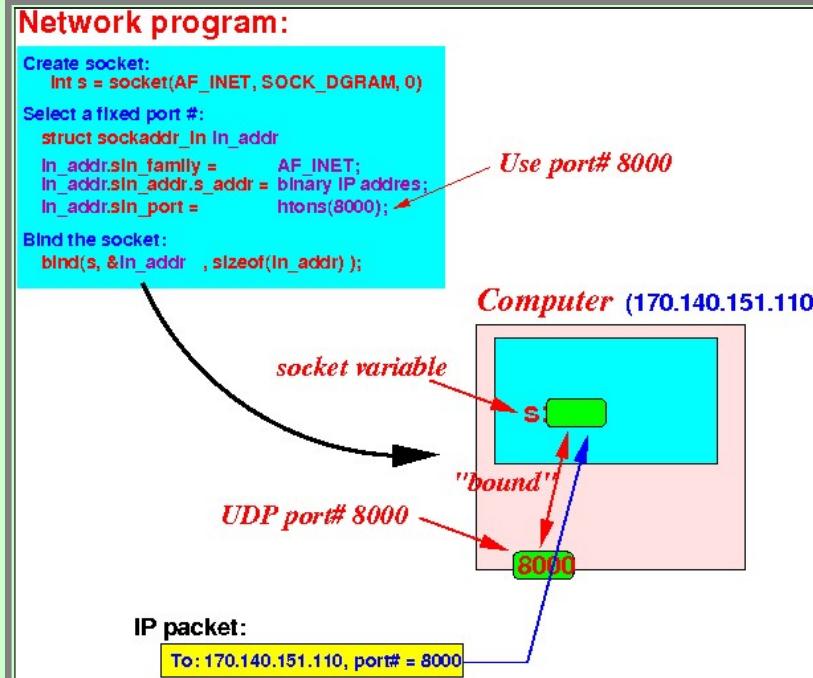
s = socket(PF_INET, SOCK_DGRAM, 0);           // Create UDP socket

/* -----
   Set up Internet network address
----- */
in_addr.sin_family = AF_INET;                /* IP Protocol domain */
in_addr.sin_addr.s_addr = htonl(IPAddress); /* IP address */
in_addr.sin_port = htons(PortNumber);        /* UDP port # */

/* -----
   Bind socket to network address
----- */
if ( bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr)) != 0 )
    perror("Bind error: " );
```

What does the program achieve:

- The following diagram shows you what the **program statements** has achieved:



Explanation:

- The **socket variable** `s` is now *related* to the **UDP port #8000** of the **computer**

- **That means:** IP messages that have **destination address**:

```
IP address = the IP address of this computer
port #     = 8000
```

will be *delivered* to the **socket s** !!!

- Example Program: (Demo above code)

Example

- Prog file: [click here](#)

Notes:

- Compile it with: **cc -o udp1 udp1.c**
- Run it with: **udp1 IP-address Port#**

Experiment:

- Try run udp1 using an IP-address that is NOT assigned to the host and see what happens
- Use the **"host"** command to find the **IP-address** of the host and run udp1 with a legal IP-address
- Also, try to bind to port# lower than 1000 (system reserved ports)

Example: (on W301 computer):

```
/*
-----*
Find W301's IP address in dotted decimal notation
-----*/
cheung@W301(1007)> ifconfig
eth0      Link encap:Ethernet HWaddr 3C:D9:2B:76:D7:09
          inet addr:170.140.151.25 Bcast:170.140.151.255 Mask:255.255.254.0
          inet6 addr: fe80::3ed9:2bff:fe76:d709/64 Scope:Link

/*
-----*
Convert IP address in binary
-----*/
cheung@W301(1008)> bc
170*256*256*256 + 140*256*256 + 151*256 + 25
2861340441

/*
-----*
Bind call with W301's IP address
-----*/
cheung@W301(1011)> udp1 2861340441 8000
**** Bind was successful....
Socket is bind to:
  addr = 2861340441
  port = 8000

/*
-----*
ERROR test: Bind call with not-your-own IP address
-----*/
cheung@W301(1010)> udp1 2861340442 8000
Error: bind to addr = 2861340442,, port = 8000 FAILED
Bind ?: Cannot assign requested address
```

- Binding to a **used** port #

- Warning:

- When a **port#** is **used** (= bound) by some **other** network program, then:

- The **bind()** call to the **used port#** will **fail** (and returns an **error code**)

- Comment

- Fact:

- There is **no need** to **bind** to a **specific port#** if:

- just want to **transmit messages**

- You **only** need to **bind** to a **specific port#** if:

- **Some other process** needs to **contact you**

I.e.:

- When you are **writing** a **server** !!!

Binding to an available (kernel) UDP port

- "System Selected" (= arbitrary selected) Port Numbers

- Fact:

■ **Most network applications** can use an *arbitrary* (available) port #

- **System selected** port number:

■ **System selected** port number = an *(unused)* port number that is **picked (selected)** by the **Operating System**

■ A **system selected** port number is just an *arbitrary (unused)* port #

- Binding to an arbitrary available (kernel) port

- How to bind to a **System Selected** (kernel) port

■ Specify **0 (zero)** for the **sin_port** variable in the **network address**

Example:

```
in_addr.sin_family = AF_INET;
in_addr.sin_addr.s_addr = htonl(INADDR_ANY);
in_addr.sin_port = htons( 0 ); /* Let system assign port # */
                                ^^^

if ( bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr)) != 0 )
    perror("Bind: ");
```

Example

- Example Program: (Demo above code)

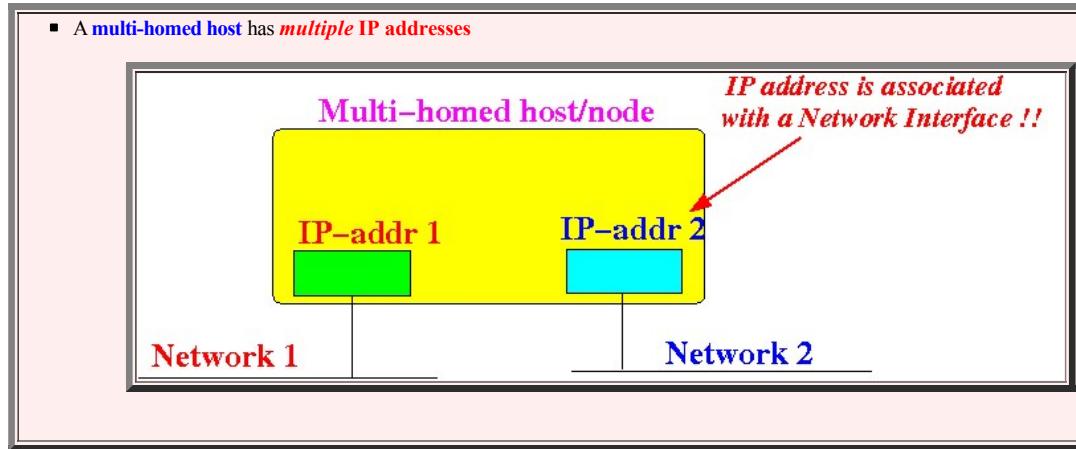
■ Prog file: /home/cs455001/demo/netw-prog-ipv4/udp3.c

Use **port = 0** in the **arguments**

Binding a UDP socket in a multi-homed host

- Multi-homed hosts

- Recall that:



- The association of a socket to a network address

- Recall the structure of a **network address**:

```
struct sockaddr_in
{
    sa_family_t      sin_family;    /* Address family
                                         - Must contain the value AF_INET */

    in_port_t        sin_port;     /* Internet Port #
                                         - a 2 bytes port number */

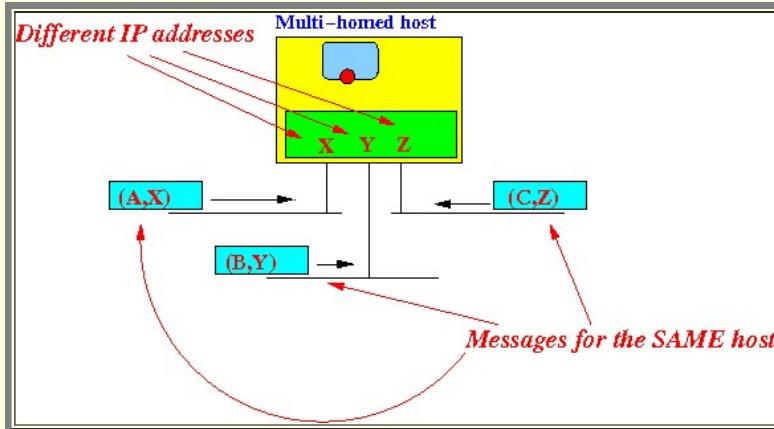
    struct in_addr sin_addr;     /* Internet Address (IP-address)
                                         - a 4 bytes IP address */
};
```

Facts:

- An **network address** contains **one IP address** value
- A **socket** that is **bound** to an **IP address X** will:
 - **Only receive IP messages** with the **destination IP address X**

- Situation:

- A **multi-homed node** can **receive** messages with **different IP addresses**:

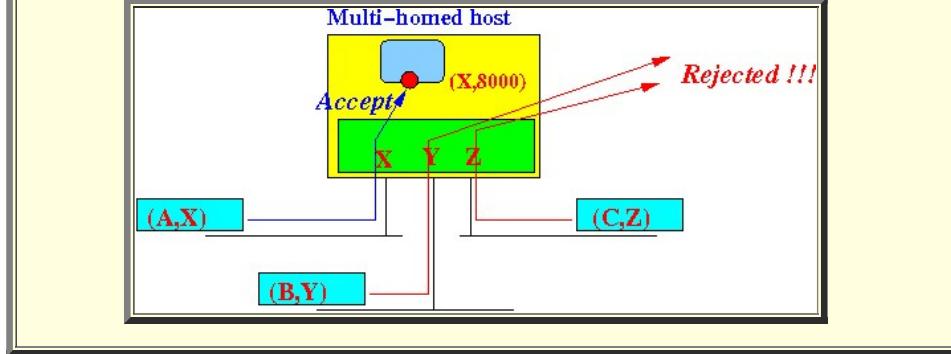


IP packets with **destination IP address**:

- **X or Y or Z**

will be **delivered** to the **multi-homed host** !!!

- A **socket** bound to a **specific IP address** will **only receive IP packets** destined for the **specific IP address**:



- Wild Card IP address

 - Wild card IP address

 - `INADDR_ANY` = the *wild card IP address* used in *network programming*

A wild card IP address will *match* to *any* IP address

 - Fact:

 - A socket that is **bound** to the `INADDR_ANY` is **able** to:

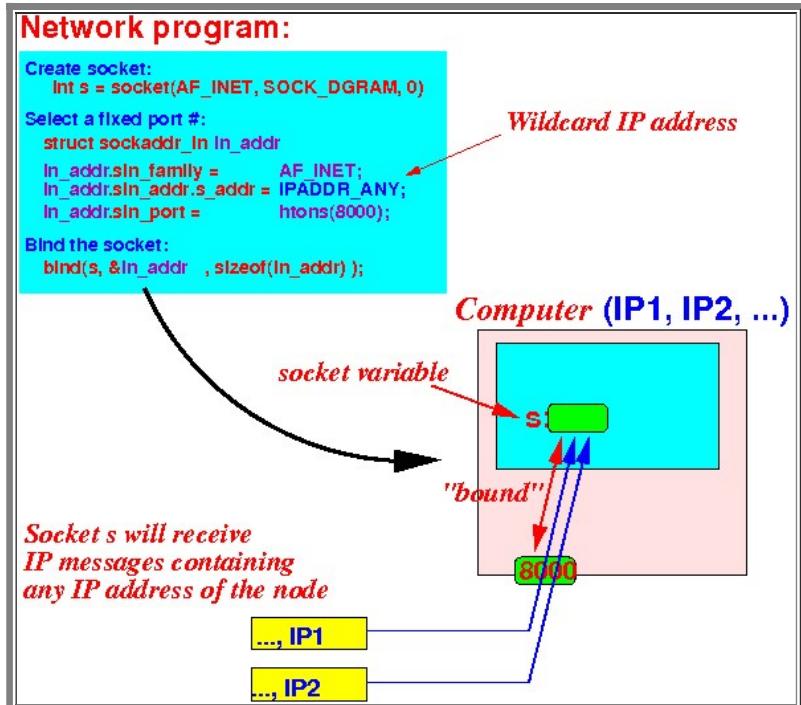
 - Receive IP packets containing *any* IP address !!!!

Example:

```
in_addr.sin_family = AF_INET; /* Protocol domain */
in_addr.sin_addr.s_addr = htonl(INADDR_ANY); /* Use wild card IP address */
in_addr.sin_port = htons(PortNumber); /* UDP port # of socket */

/* -----
   Bind socket to socket address
----- */
if ( bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr)) != 0 )
    perror("Bind error: ");
```

Result:



- Common practice of network programming

 - Common practice in *network programming*:

 - Network programmers *always* use the `INADDR_ANY` value in an *network address*

Example:

```
in_addr.sin_family = AF_INET;
in_addr.sin_addr.s_addr = htonl(INADDR_ANY);
/* Always use wild card IP address */
in_addr.sin_port = htons(PortNumber);

/* -----
   Bind socket to socket address
----- */
if ( bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr)) != 0 )
    perror("Bind error: ");
```

- It is a **general practice** to use **INADDR_ANY even** for **single-homed** computers...
(Because it's **easier**)

Example

- Example Program: (Bind to **INADDR_ANY**)

- Prog file: [click here](#)

Notes:

- Compile it with: **gcc -o udp2 udp2.c**
- Run it with: **udp2 Port#**

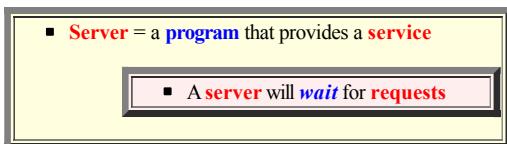
Example:

```
cheung@compute(1011)> udp2 8000
OK: bind to port = 9000 SUCCESS
**** Bind was successful....
Socket is bind to:
    addr = 0
    port = 8000
```

Client/server programming with UDP sockets

- Server

- Server:



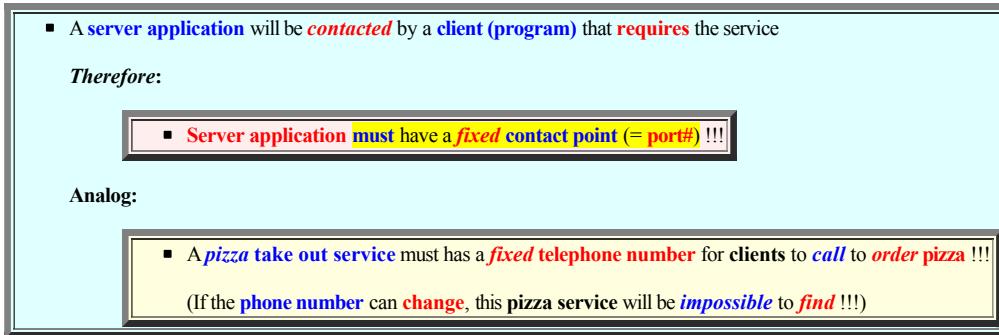
- Client

- Client:



- Writing a server program

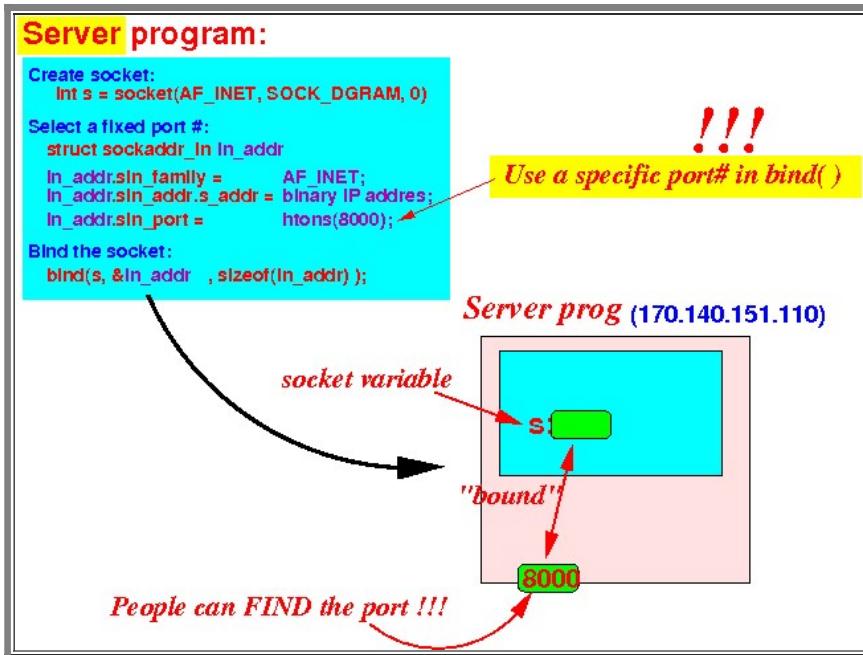
- Server programming:



Analog:

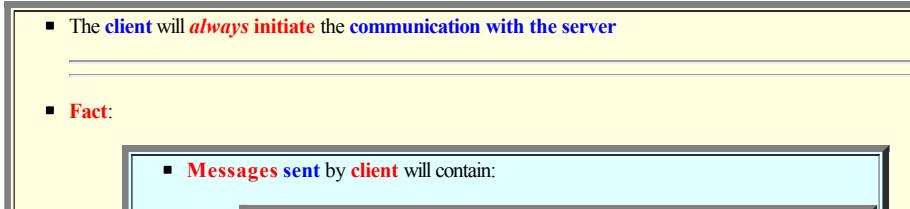
- A **pizza take out service** must has a **fixed telephone number** for clients to **call** to **order pizza !!!**
(If the **phone number** can **change**, this **pizza service** will be **impossible** to **find !!!**)

- Summary: when **writing a server**, use a **fixed (UDP) port#**



- Writing a client program

- Client programming:



■ The IP address and the port# of the socket of the client program !!!

Analogy:

■ We **you** calls to **order pizza**, the **pizzaria** will have your **phone number** (as in **caller ID**)

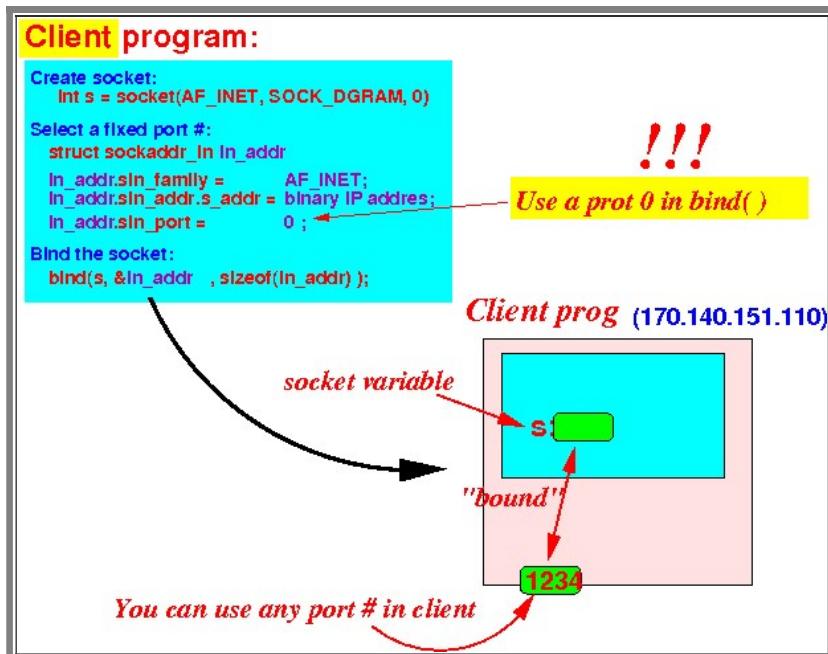
■ Consequently:

- The **server** can **find out** the **network address** of the **client** !!!
- A **client** can **use a arbitrary (available) port number** !!!

◦ Conclusion:

■ When **writing** a **client application**, **bind** the **socket** to a **system assigned port #**

◦ Summary



Sending and receiving with UDP sockets

- Sending Packets using a UDP port

- System call to **send** data with (bound) UDP socket:

```
int sendto(int sock,  
          void *data,    int len,           // Transmitted data  
          int flags,      // Transmit options  
          struct sockaddr *to,   int to_len); // Destination network address
```

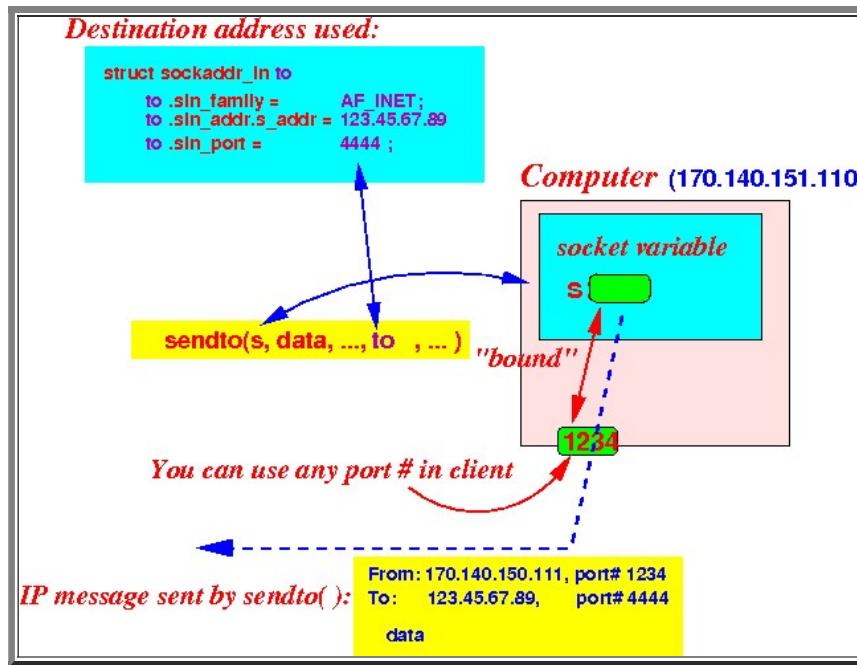
Meaning of the parameters:

```
sock = socket variable (should be bounded)  
data = pointer to data buffer location  
len = number bytes in data buffer to be sent  
flag = flags (indicate send options)  
       E.g., Set MSG_OOB bit for "out-of-band" (high priority) data  
to = Internet address (typed) variable  
     contains: network address of the destination  
to_len = length of the "to" data structure
```

- Effect:

- Transmits a UDP message containing the **data** to the specify **destination** (IP address + UDP port#)
 - The UDP mesasge will contain the **network address** (IP addr + UDP port#) of the **socket s**

Graphically:



- Return value:

```
Error:      returns -1  
Otherwise:  returns len (>= 0)
```

General error processing:

- If a **system call** (such as **sendto** returns an **error**), you can use:

```
perror( "Notice text" ); // print error
```

to **print** the **error message** to the **terminal**

- Example: send the string "Hello"

```
char data[] = "Hello"; // String has 6 characters: H e l l o \0 !!!  
struct sockaddr_in dest_addr; // Internet address type
```

```

/* -----
   Set up destination: (IPaddress, Port#)
----- */

dest_addr.sin_family      = AF_INET;
dest_addr.sin_addr.s_addr = htonl(DestIPaddr);
dest_addr.sin_port         = htons(DestPortNumber);

if ( sendto(s, data, 6,
             0 /* no flags */,
             (struct sockaddr *)&dest_addr, sizeof(dest_addr))
    == -1 )
{
    perror("sendto()");
}

```

- Example Program: (Demo above code)

Example

- Prog file: [click here](#)

How to run the program:

- Right click on link(s) and save in a scratch directory
- To compile: `gcc -o udp4-s1 udp4-s1.c`
- To run: `udp4-s1 Dest-IP-addr Dest-Port#`

Example: `udp4-s1 2827975501 (=170.140.150.1) 9000`

Note:

- use **host NameOfComputer** and **bc** to compute the destination IP address (in binary)

- **Warning:** you won't see anything....

- because the **receiver** is **not yet running** :-)

• Remarks about the UDP send program

- Important observation:

- When you **run** the **above demo program**, you should **notice that**:

- The program can send a **UDP message** even when there is **no receiver !!!**

This can be done **because**:

- **UDP messages** are **unreliable**
- I.e.: `sendto()` does **not** require any **acknowledgement** from a **receiver !!!**

Analogy:

- `sendto()` is like **mailing a letter**
- The **sender** does **not** know the **result** of the **transmission**

• Sending with an **unbounded** UDP socket

- Hacker's note:

- It is **possible** to send **UDP messages** using an **unbounded UDP socket**:

```

int    IPAddress;           /* a 32 bit IP address */
short  PortNumber;          /* a 16 bit port number */

int    s;
struct sockaddr_in in_addr;

s = socket(AF_INET, SOCK_DGRAM, 0);           /* Create socket
// s is unbounded

/* -----
   OMIT the bind step
----- */
in_addr.sin_family = AF_INET;                /* Protocol domain */
in_addr.sin_addr.s_addr = htonl(IPAddress); /* IP address of socket */

```

```

in_addr.sin_port = htons(PortNumber); /* UDP port # of socket */

/*
----- Bind socket to socket address -----
*/
if ( bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr)) != 0 )
    perror("Bind error: ");

/* -----
   Set up destination: (IPaddress, Port#)
----- */
dest_addr.sin_family      = AF_INET;
dest_addr.sin_addr.s_addr = htonl(DestIPAddr);
dest_addr.sin_port        = htons(DestPosrtNumber);

if ( sendto(s, data, 6,
            0 /* no flags */,
            (struct sockaddr *)&dest_addr, sizeof(dest_addr))
    == -1 )
{
    perror("sendto()");
}

```

(An **unbounded socket** is a **socket variable s** where you has **not** called **bind(s, ...)**)

- o Effect of sending with an unbounded socket:

- UDP Messages **sent** with an **unbounded UDP socket** is:

■ "undefined"
(= depends on the UDP implementation)

- Possible results:

■ Some system will bind the socket (to an available UDP port) first

■ Other systems will send UDP messages with:

■ Empty (= garbage) sender's IP address and port number
--

(I.e.: anonymous IP messages !!!)
--

Receiving with UDP sockets

- Receiving Packets using UDP port

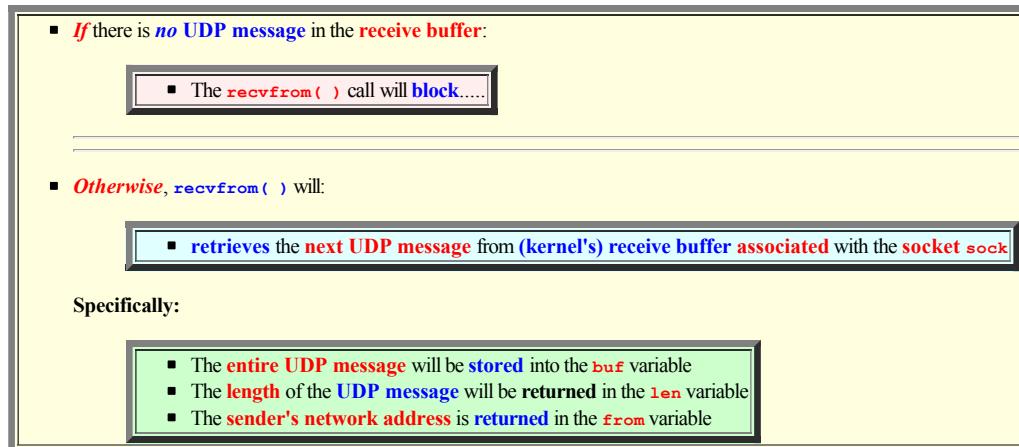
- System call to receive data with UDP socket:

```
int recvfrom(int sock,
             void *buf, int len,           // Receive buffer
             int flags,
             struct sockaddr *from, int *from_len); // Sender's network addr.
```

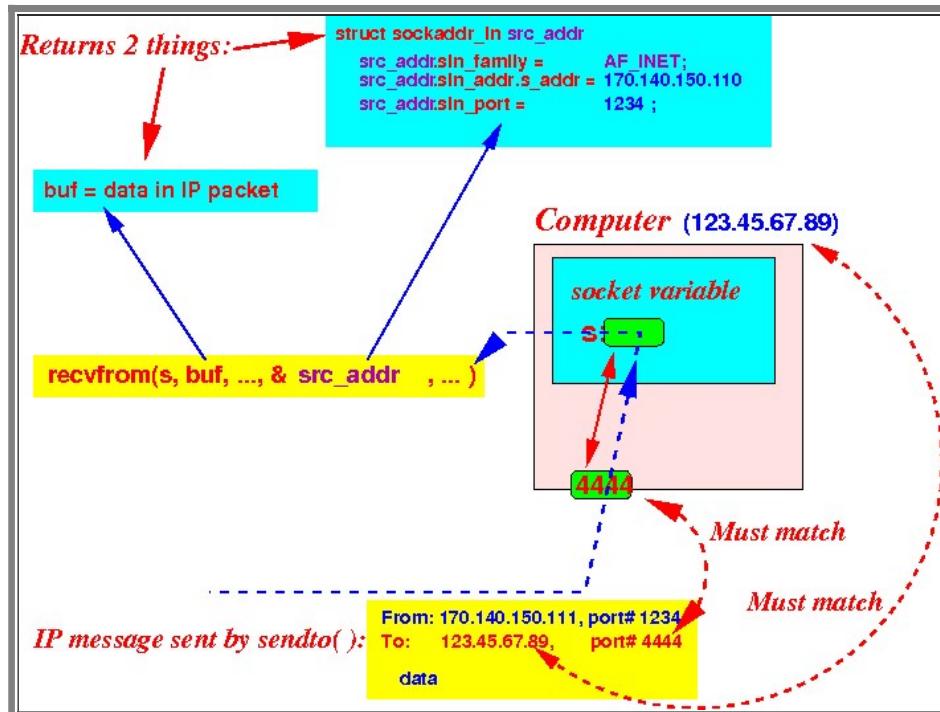
PARAMETER:

- sock = socket variable
- buf = pointer to user data buffer location (stores recv'd message)
- len = buffer size of "buf"
- flag = various flags (options)
E.g., Set MSG_OOB bit for "out-of-band" (high priority) data
- from = IP address + Port number of the sender !!!
- from_len = length of the "from" structure

- Effect:



Graphically:



- Very important fact:

- The **network address** (= IP addr + port#) of the **sender** is **return** in:
 - The variable **from**

Therefore:

- The receiver can send back a reply using:

```
sendto( s, reply, len, 0, from, from_len);
```

- Return value: (of `recvfrom()`)

```
Error: returns -1  
Otherwise: returns the number bytes stored in buf
```

Example:

```
/* =====  
 Assume we have done these already:  
  
     create socket s  
     bound s to a port  
===== */  
  
char line[1000];           /* User buffer to receive data */  
  
struct sockaddr_in  src_addr;      /* Used to store sender's address */  
int    length;                /* Length of the sockaddr_in data type */  
  
/* -----  
 Set up the length to tell recvfrom how much space  
 it can write into the "src_addr" variable  
----- */  
length = sizeof(src_addr);        /* Length of the src_addr variable */  
  
/* -----  
 Receive the next message into the variable "msg"  
 The sender of the message will be recorded in "src_addr"  
 (You can use "src_addr" to reply to the sender !)  
----- */  
recvfrom( s, line, 1000, 0 /* flags */,  
         (struct sockaddr *)&src_addr, &length);
```

Example

- Example Program: (Demo above code)

- Prog file: [click here](#)

How to run the program:

- Right click on link(s) and save in a scratch directory
- To compile: `cc -o udp4-rl udp4-rl.c`
- To run: `udp4-rl My-IP-address my-Port#`
 - use **My-IP-address = 0** for **wildcard IP address**
 - You need to use the **same port#** in the **send program above** to send message to this program

Sending a *reply* back to the sender

- You can **send a reply** back to the **sender** using the `src_addr` variable from the `recv_from()` call:

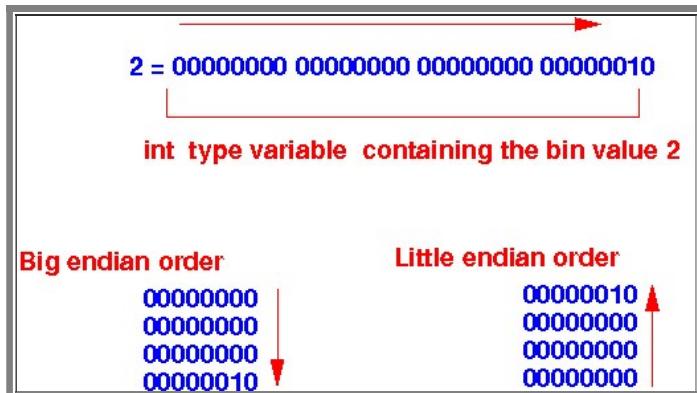
```
/* -----  
 Receive the next message into the variable "msg"  
 The sender of the message will be recorded in "src_addr"  
 (You can use "src_addr" to reply to the sender !)  
----- */  
recvfrom( s, line, 1000, 0 /* flags */,  
         (struct sockaddr *)&src_addr, &length);  
  
...  
  
/* -----  
 Program can send reply to sender using "src_addr"  
----- */  
sendto( s, buf, length_of_buf, flags,  
       (struct sockaddr *)&src_addr, &length);
```

Sending and receiving binary numbers

- Storing integer (or any other type of) binary numbers
 - Computers store binary numbers in one of 2 possible order:

- Big endian ordering
- Little endian ordering

Graphically:



- Host byte ordering:

- Host byte ordering = the storage ordering of binary numbers used in a computer (= host)

- Transmitting integer (or any other type of) binary numbers

- Fact:

- The function call:

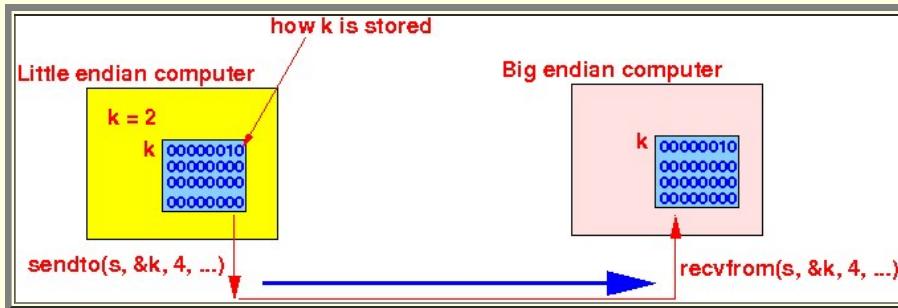
```
sendto( s, data, len, flag, toAddr, toAddrLen )
```

will transmit the `len` bytes (stored) in the variable `data` in the `order` stored in (host) memory

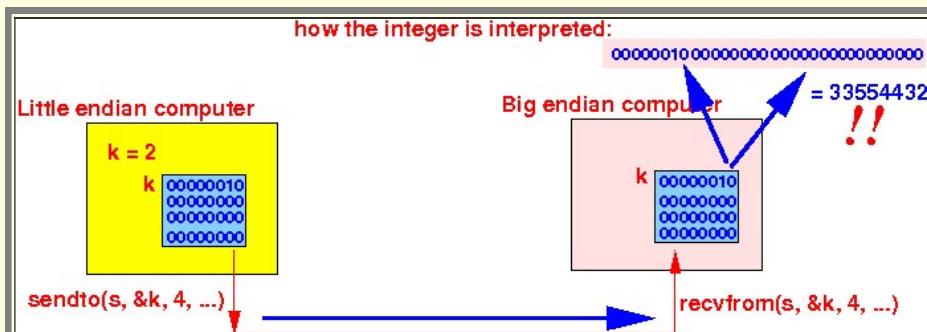
- What happens if we transmit binary data between a big endian computer and a little endian computer

- Consider the following data transmission of an integer binary number from little endian computer to a big endian computer:

- A little endian computer send the integer value 2 to a big endian computer



- The big endian computer will interpret the integer binary number as follows:



- Example Program: (Demo above code)

Example

- Prog file:

- /home/cs455001/demo/netw-prog-ipv4/udp-send-int.c
- /home/cs455001/demo/netw-prog-ipv4/udp-recv-int.c

How to run the program:

- To compile:

- gcc -o udp-send-int udp-send-int.c on W301 (or a lab machine)
- gcc -o udp-recv-int udp-recv-int.c -lnsl -lsocket on compute

- To run:

- On compute: udp-recv-int 8000
- On W301: udp-send-int compute 8000

- How to transmit binary data

- Technique to transmit binary data:

- Sender:

1. First, convert the binary data to **network byte ordering**

2. Then, send the **converted data**

Example:

```
int k = 1;      // k is stored in host byte order
printf("Sending: int value = %d\n", k);
int x = htonl(k);      // Convert to network byte ordering
sendto(s, &x, sizeof(int), 0 /* flags */,
       (struct sockaddr *)&dest_addr, sizeof(dest_addr));
```

- Receiver:

- Receiver the **binary data** (we know it is in **network byte order**)

- Convert the data to **host byte order** using:

```
int ntohs( int ): convert 4 byte data from network byte ordering
                  to host byte ordering

short htons( short): convert 2 byte data from network byte ordering
                     to host byte ordering
```

Example:

```
int x;
src_addr_len = sizeof(src_addr);
len = recvfrom(s, &x, sizeof(int), 0 /* flags */,
               (struct sockaddr *) &src_addr, &src_addr_len);
k = ntohs( x );      // Convert int from network byte ordering
                      // to host byte ordering
printf("k = %d", k);
```

- Example Program: (Demo above code)

Example

- Prog file:

- Edit the same demo programs:

■ /home/cs455001/demo/netw-prog-ipv4/udp-send-int.c
■ /home/cs455001/demo/netw-prog-ipv4/udp-recv-int.c

Uncomment the `htonl()` and `ntohl()` calls in the code

How to run the program:

- To compile:

■ gcc -o udp-send-int udp-send-int.c on W301 (or a lab machine)
■ gcc -o udp-recv-int udp-recv-int.c -lnsl -lsocket on compute

- To run:

■ On compute: udp-recv-int 8000
■ On W301: udp-send-int compute 8000

• Transmitting structured data

◦ Fact:

- Structured data can contain:

■ String/text data (stored in the <i>same order</i> in both big endian computer and little endian computer)
■ integer/float binary data (stored in <i>different order</i> in big endian computer and little endian computer)

◦ How to transmit *structured* data:

- Transmitting the *structured* data:

1. Use `htonl()` (or `htons()`) to convert every integer/float variable to network byte order
2. Send the converted struct using `sendto()`

◦ How to receive *structured* data:

- Transmitting the *structured* data:

1. Use `recvfrom()` to receive the transmitted struct data
2. Use `ntohl()` (or `ntohs()`) to convert every integer/float variable to host byte order

Example

◦ Example Program: (Demo above code)

- Prog file:

■ /home/cs455001/demo/netw-prog-ipv4/udp-send-struct.c
■ /home/cs455001/demo/netw-prog-ipv4/udp-recv-struct.c

How to run the program:

- To compile:

■ gcc -o udp-send-struct udp-send-struct.c on W301 (or a lab machine)
■ gcc -o udp-recv-struct udp-recv-struct.c -lnsl -lsocket on compute

- To run:

■ On compute: udp-recv-struct 8000
■ On W301: udp-send-struct compute 8000

Intro to "Support Functions" for Network Programming

- Network Support Functions

- Fact:

- Computer on the Internet are identified by their IP addresses

- When we wish to visit google.com, we need to use its IP address !!!

(Analogy: dial the telephone number)

- IP-addresses are binary numbers

- It is tedious to remember a binary number !!!

- Functions to support network programming:

- Lookup function to find the corresponding IP address of a symbolic host name

Example:

- You are given the symbolic host name, e.g.: www.google.com

- Find the IP address for www.google.com

Analogy: look up in a telephone directory

- Reverse look up function:

- You are given an IP address

- Find the symbolic host name for that IP address

- Dotted decimal address to binary IP address translation:

- You are given a dotted decimal IP address, e.g.: "170.140.150.1"

- Find the corresponding binary IP address

- The network service library

- The network service library for C:

- network service library (nsl) = a function library in C programming language that contains the IP address conversion functions

- Compiling programs that uses the network service library functions:

- When your network program uses some network service function(s):

- You need to link the network service library (nsl) when you compile your program

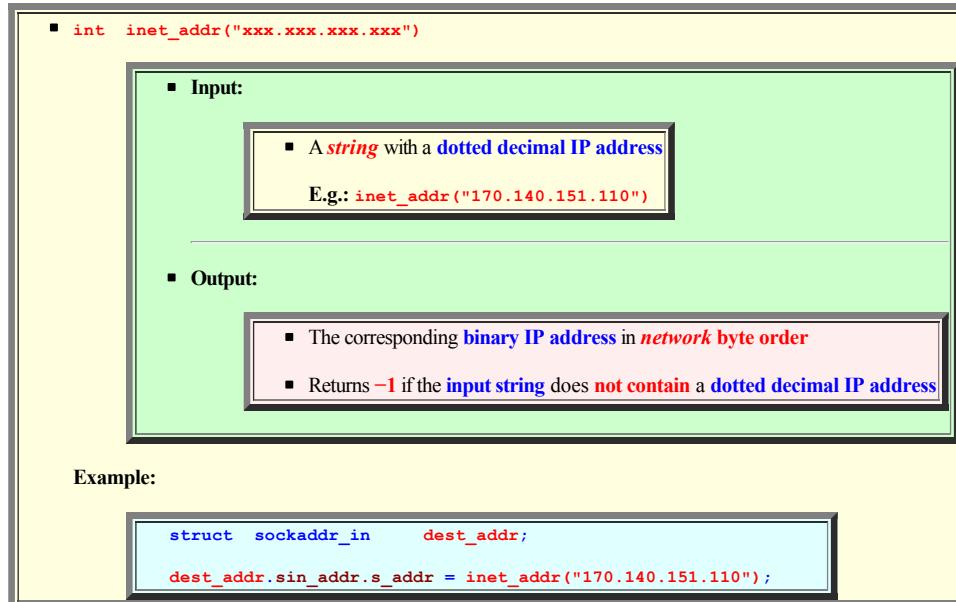
- How to link the network service library:

```
cc -o my-netw-prog my-netw-prog.c -lnsl
```

Converting "dotted-decimal-format" to (binary) IP address

- "Dotted Decimal Notations" ==> *binary* IP addresses conversion

- Converting a "Dotted Decimal Notation" string (such as "170.140.151.110") to a **binary IP address**:



- **Important note:**

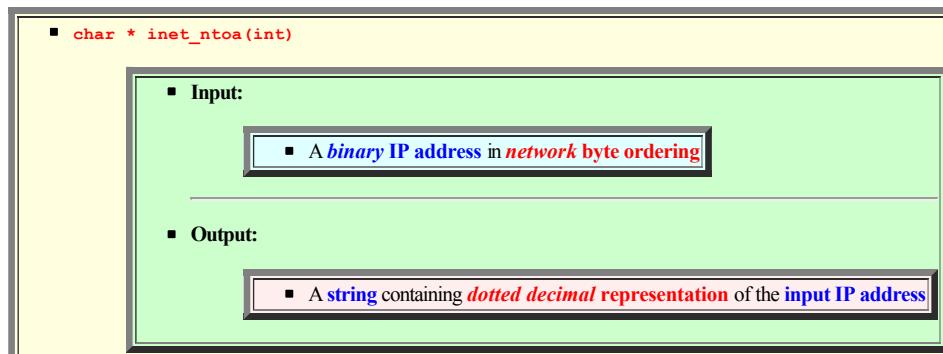
- **Do not** use `htonl()` on the **output** of `inet_addr()`

Because:

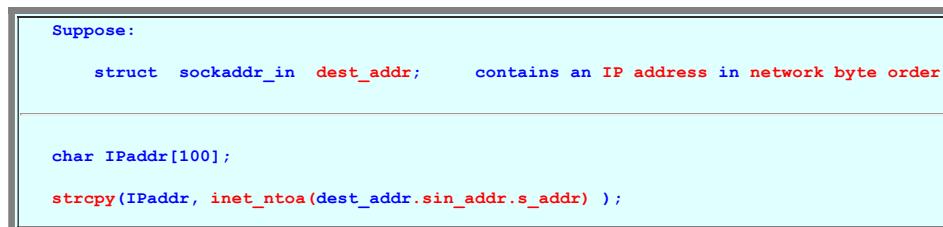
- the **output** of `inet_addr()` is **already** in **network byte ordering !!!**

- (Binary) IP address ==> "dotted-decimal" conversion

- Converting a **binary IP address** *back* to a "Dotted Decimal Notation" string (such as "170.140.151.110"):

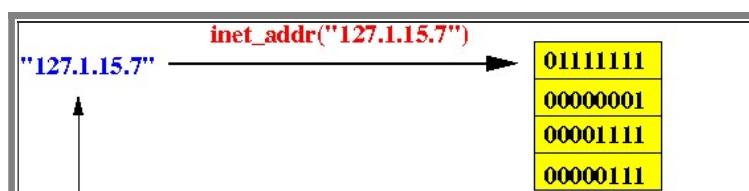


Example:



- **Summary:** `inet_addr` and `inet_ntoa`

- The **following figure** summarizes the **operation** of the `inet_addr()` and `inet_ntoa()` functions:



inet_ntoa(int)

- Network programming using *dotted decimal* IP address

- The **UDP sender** using *dotted decimal* IP address to specify the **receiver**:

```
int main(int argc, char **argv)
{
    struct sockaddr_in dest_addr;           /* Destination socket address */

    ...

    /* =====
       argv[1] is the string containing the dotted decimal IP address
    ===== */
    if ( ( dest_addr.sin_addr.s_addr = inet_addr( argv[1] ) ) == -1 )
    {
        printf("Dotted IP-address must be of the form a.b.c.d\n");
        exit(1);
    }

    ...
}
```

- Example Program: (Demo above code)

Example

- **Sender:** /home/cs455001/demo/netw-prog-ipv4/udp4-s2.c
- **Receiver:** /home/cs455001/demo/netw-prog-ipv4/udp4-r2.c

- The new UDP sender using *dotted decimal* IP address Prog file: [click here](#)
- The (**unmodified**) UDP receiver Prog file: [click here](#)

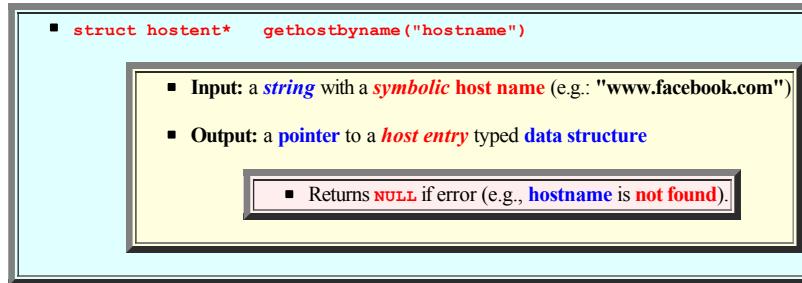
How to run the program:

- To compile the **sender**: `gcc -o udp4-s2 udp4-s2.c -lns1`
 - To compile the **receiver**: `gcc -o udp4-r2 udp4-r2.c`
-
- To run the **receiver**: `udp4-r2 port#`
 - To run the **sender**: `udp4-s2 xx.xx.xx.xx port#`

Converting "symbolic names" to (binary) IP address

- `gethostbyname()`: Find *host information* using a *Symbolic Host name*

- Converting a *symbolic host name* (such as "www.facebook.com") to a *binary IP address*:



- The "host entry" data structure

- The *host entry* `struct hostent` data structure:

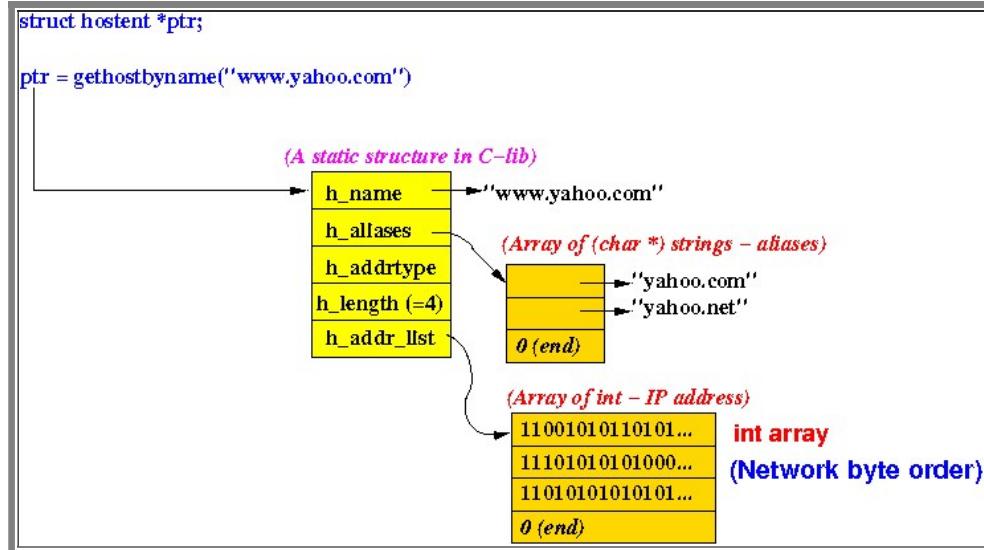
```
struct hostent
{
    char * h_name;          /* The canonical (official) name of host */
                           // a string variable

    char * h_aliases[ ];   /* The list of "aliases" for the hostname */
                           // h_aliases is an array of string
                           // h_aliases[i] == NULL marks the end of the array

    int    h_addrtype;     /* host address type (should be = AF_INET (2)) */
    int    h_length;        /* length of address (= 4 bytes IP addresses) */

    char * h_addr_list[ ]; /* list of addresses */
                           // h_addr_list is an array of int (binary IP addresses)
                           // in network byte order
                           // h_addr_list[i] == 0 marks the end of the array
};
```

- Graphical representation of the "host entry" data structure (with *sample content* to make things clearer):



- How to use/access the data in a host entry (too much info)

Example: how to use the `gethostbyname()` call

```
/* -----
   Variables
----- */
char hname[100];
struct hostent *hp; /* host pointer */

int addr;
int i;

printf("Enter hostname: ");
scanf("%s", hname);

hp = gethostbyname(hname);

if (hp == NULL)
```

```

    printf("host information for %s not found\n", hname);
    exit(1);
}

/* -----
   Print the information
----- */

/* -----
   Canonical name
----- */
printf("official hostname h_name = `%s'\n", hp->h_name);

/* -----
   List of aliases
----- */
printf("Host has this alias list:\n");

char *p[];           /* help variable to access h_aliases[ ] */

p = hp->h_aliases; // p[i] is a string (char *)

for ( i = 0 ; p[i] != NULL; i++ )
{
    printf("\tAlias %d: %s\n", i+1, p[i] );
}

printf("Host's address type = %d (should be AF_INET = %d)\n",
       hp->h_addrtype, AF_INET);

/* -----
   Length of an IP address
----- */
printf("Host's address length = %d\n", hp->h_length);

/* -----
   List of IP addresses
----- */
printf("List of Internet addresses for host:\n");

int *a[];           /* Help variable to access h_addr_list[ ] */

a = (int **) hp->h_addr_list; // a[i] is an IP address

for (i = 0; a[i] != 0; i++)
{
    printf("\tIP address %d: ", i+1);
    paddr( (void*) a[i] );      // Print IP address as dotted decimal
    printf("\n");
}

```

- Example output:

```

Enter hostname: www.yahoo.com

official hostname h_name = `ds-any-fp3-real.wal.b.yahoo.com'

Host has this alias list:
    Alias 1: www.yahoo.com
    Alias 2: fd-fp3.wg1.b.yahoo.com
    Alias 3: ds-fp3.wg1.b.yahoo.com
    Alias 4: ds-any-fp3-lfb.wal.b.yahoo.com

Host's address type = 2 (should be AF_INET = 2)

Host's address length = 4

List of Internet addresses for host:
    IP address 1: 98.138.253.109
    IP address 2: 206.190.36.45
    IP address 3: 98.138.252.30
    IP address 4: 206.190.36.105

```

- Example Program: (Demo above code)

Example

- Prog file: [click here](#)

Notes:

- Compile with: cc -o gethostbyname gethostbyname.c -lssl

- Network programming using symbolic host names

- Fact:

- We just need a (any) IP address of a IP host !!!

- Variable in the host entry structure that contain the IP addresses of the host:

```
h_addr_list[ ]      // It's an array (multiple IP addresses)
```

We just need 1, so the safest choice is:

```
h_addr_list[0]    // The other entries may be invalid !!!
```

- The UDP sender using a *symbolic host name*:

```
int main(int argc, char **argv)
{
    struct sockaddr_in dest_addr;           /* Destination socket address */
    struct hostent *hp;

    ...

    /* =====
       argv[1] is the string containing the symbolic host name
    ===== */
    hp = gethostbyname(argv[1]);           // E.g. argv[1] = "www.google.com"

    if ( hp == NULL )
    {
        printf("Error... \n");
        exit(1);
    }

    /* =====
       Use:
       hp->haddr_list[0]      // This one is for sure valid :)

       Note: hp->haddr_list[0] is in network byte ordering !!!
    ===== */

    memcpy( (char *) &(dest_addr.sin_addr.s_addr),
            hp->h_addr_list[0],          // Copy first entry and
            hp->h_length);             // preserve byte ordering

    ...
}
```

Note:

- `memcpy(dest, src, len)` copies `len` bytes from `src` address to `dest` address while *preserving the byte ordering*
 - **Do not** use `htonl()`

*because the IP address in `h_addr_list[0]` is *already* in network byte ordering !!!)*

- Example Program: (Demo above code)

Example

- The new UDP sender using *symbolic host names* Prog file: [click here](#)
- The (*unmodified*) UDP receiver Prog file: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory
- To compile the **sender**: `gcc -o udp4-s3 udp4-s3.c -lnsl`
- To compile the **receiver**: `gcc -o udp4-r2 udp4-r2.c`
- To run the **receiver**: `udp4-r2 port#`
- To run the **sender**: `udp4-s3 lab???.mathcs.emory.edu port#`

Find host information using an IP address

- **gethostbyaddr()**: Find *host information* from an IP address

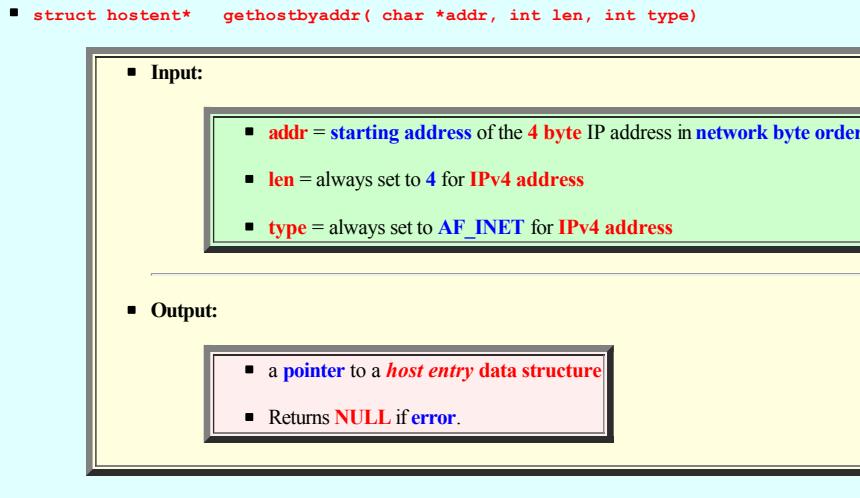
- This **functionality** is *rarely used*:

- Find **host information** (e.g., the *symbolic host name*) using a *binary IP address*

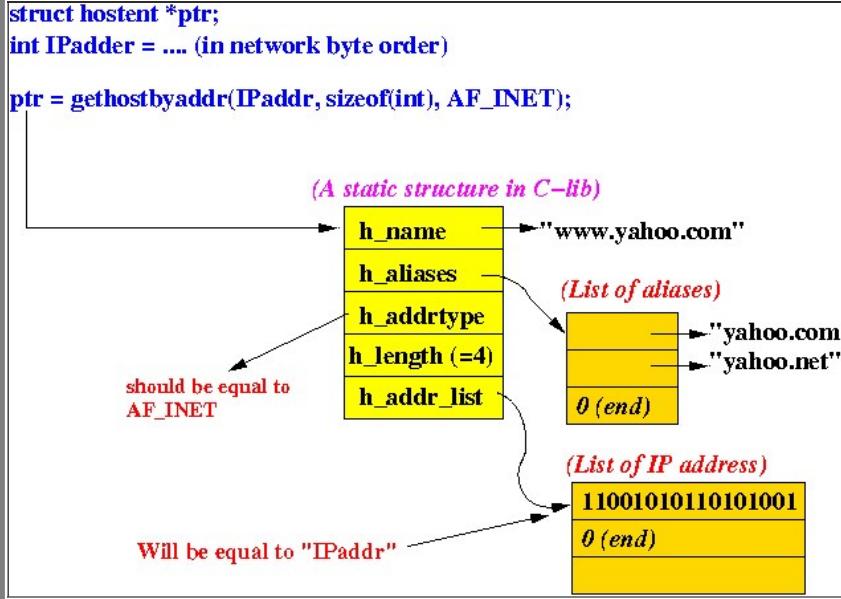
In case you **do** need it, the **function** is:

```
gethostbyaddr( ... )
```

- Find **host information** using its *binary IP address*:



- `gethostbyaddr()` will *also* return a "host entry" data structure (just like `gethostbyname()`):



(It is the **same data structure** returned by `gethostbyname()`)

Therefore, I will **skip** on how to use this data structure....

You can read the **program** below yourself....)

- **Example Program:** (Demo above code)

Example

- Prog file: [click here](#)

Notes:

- Compile with: `cc -o gethostbyaddr gethostbyaddr.c -lsl`

Programming with Timeout

- Recall: the Stop-and-wait protocol

- The stop-and-wait protocol (click here [click here](#)) transmits a message and must wait for acknowledgement:

```
stop-and-wait-send( input-data )
{
    seqno++;           // use next sequence number

    msg.seqno = seqno;      // label message with the seqno
    msg.data  = input-data // put input data inside msg

    send msg;           // send message to receive (use sendto( ) !)

    ack_recv = false;

    while ( ! ack_recv )
    {
        wait for acknowledgement;

        if ( timeout while waiting for ACK )
        {
            send msg (again);      // retransmit message
        }
        else
        {
            if ( ACK.seqno = seqno )
                ack_recv = true;
        }
    }
}
```

- Functions that support "Timeout programming"

- Time out:

- Network programs often use timeout to recover from transmission errors

- System calls (= functions) that can be used in timeout programming:

- **select()** (originally defined in the **BSD (Berkeley Software Distribution)** UNIX)
- **poll()** (originally from **UNIX System V**)

Most UNIX systems nowadays provide both system calls.

Intro to the `select()` function

- Header files for `select`

- In order to use the `select()` function in C, you must include this header file:

```
#include <sys/select.h>
```

- Usage of the `select()` function

- Usage of the `select()` function:

- `select()` allow a program to **check/test** and/or **wait until**:
 - one or more **file descriptors/sockets** is "ready" for **reading (receiving)** or **writing (sending)**

- "Ready" for reading/writing:

- Recall that:
 - A **read/recvfrom** operation will **block** if:
 - There is **no data** in the **receive buffer**
 - A **socket** is **ready** for some **operation** (read/write) if:
 - The **operation** (read/write) will **not block**

- The `select` system call

- Syntax:

```
#include <sys/select.h>
#include <sys/time.h>           /* Header files */

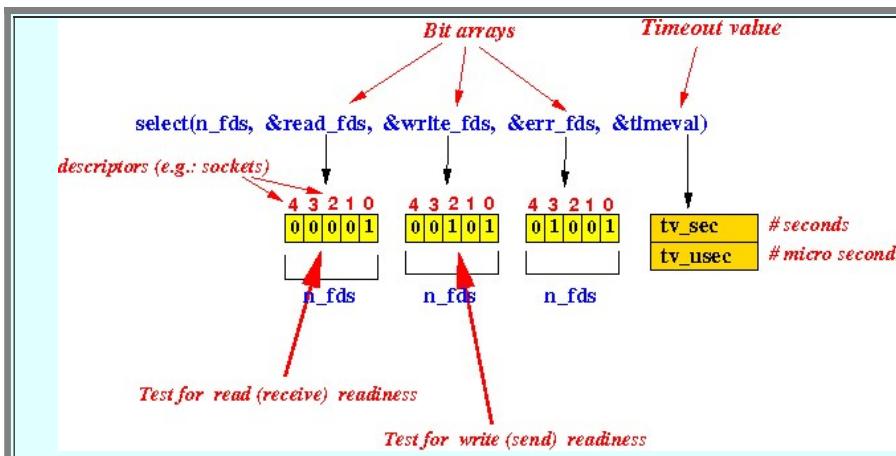
int select( int n_fds,
            fd_set * readfds,    // a bit array
            fd_set * writefds,   // a bit array
            fd_set * errorfds,  // a bit array
            struct timeval * timeout );
```

The data types used:

```
fd_set = a bit array type (an int array interpreted as bit array)

struct timeval
{
    unsigned tv_sec;      /* # seconds in the timeout value */
    unsigned tv_usec;     /* # microseconds */
}
```

Meaning of the parameters:



```

nfds      = number of bits used in readfds, writefds and errorfds

readfds   = descriptor bit array to test for read readiness
            (NULL means empty set)

        The ith bit in the bit array corresponds to the file descriptor i

writefds = descriptor bit array to test for write readiness
            (NULL means empty set)

errorfds = descriptor bit array to test for error conditions
            (NULL means empty set)

timeout   = maximum wait duration (time out value)
            (NULL means wait indefinitely)

```

◦ Effect of the `select()` function:

- The `select()` function examines the **file (or sockets) descriptor sets** in the variables:

- `readfds` for **read readiness** (= has data)
- `writefds` for **write readiness** (= has buffer space)
- `errorfds` for **error**

for a **maximum duration** given by the `timeout variable`

- If **some** of these **fds/sockets** are **ready before** the `timeout` expires:

- The `select() call` will return **immediately**

Return values:

1. The **function's return value** is **equal to**:

- the **number of file descriptors/sockets** that are **ready**

2. The **bit array** variables `readfds`, `writefds` and `errorfds` will contain:

- the **file descriptors** that are **ready** for the **corresponding operation**
(I.e., the **bit position** of the **file descriptors** in the **bit array** = **1**)

3. The **timeout** variable will be:

- **reduced** by the **amount of time** spent inside the `select() call`

(I.e.: the `timeout` variable will **contain** the **remaining waiting time** before the `timeout expires`)

- **Otherwise (nothing is ready after waiting for timeout amount of time):**

- `select()` will return (the **function value**) **0**

Simplified use of the `select()` function

- A simplified usage of `select()`

- In general:

- The `fd_set` variable used in `select()` is:

- a bit array of 1024 bits

Manipulating (set/reset) the bits in such a large bit array requires the use of macros

- However:

- you can use smaller of bit array variables in `select()` !!!!

Example:

- An `int` variable is a bit array of 32 bits

I will first show you how to use `select()` with an `int` (32 bit) bit array.

- Testing read-readiness of a socket

- Recall:

- An `int` is an array of 32 bits

- We can clear all bits in an `int` using:

- `int s = 0; // Clear all bits`

- We can set the bit position `i` using:

- `s = s | (1 << i);`

- How to test the read (receive) readiness of a socket using `select()`:

```
int fdset;          // bit array of 32 bits
int s;              // socket
// Assume s is initialized and bound to some port...

/* =====
   Set up timeout value
   ===== */
struct timeval timeout;

timeout.tv_sec = 0;
timeout.tv_usec = 50; // 50 micro sec time out

/* =====
   Prepare file descriptors for select()
   ===== */
fdset = 0;           // Clear bit array
fdset = fdset | (1 << s); // s = an initialized socket variable

/* =====
   Check for message arrival with in timeout
   ===== */
if (select(32, &fdset, NULL, NULL, &timeout) != 0)
{
    // s is read for reading
}
else
{
    // timeout occurred
}
```

- Example program on using `select()`

- Example program to demonstrate the effect of `select()`:

```

// I need to create a socket and bind it first.....

/* ---
   Create a socket
--- */
s = socket(PF_INET, SOCK_DGRAM, 0);

/* ---
   Set up socket end-point info for binding
--- */
in_addr.sin_family = AF_INET;           /* Protocol family */
in_addr.sin_addr.s_addr = htonl(INADDR_ANY); /* Wildcard IP address */
in_addr.sin_port = htons(atoi(argv[1]));  /* Use any UDP port */

/* ---
   Bind socket s
--- */
bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr))

/* =====
   How to use select
===== */
int recv_fds;                      /* I can use a simple int */
struct timeval timeout;            /* Time value for time out */

/* =====
   Select the bit for socket s
===== */
recv_fds = 0;                      // Clear all bits in recv_fds
recv_fds = recv_fds | (1 << s); // Set bit for socket s

/* =====
   Set timeout value to 5 sec
===== */
timeout.tv_sec = 5;                // Timeout = 5 sec + 0 micro sec
timeout.tv_usec = 0;               // number of micro seconds

int result;

printf("Now send a UDP message to port %s before %d sec !!!...\n",
       argv[1], timeout.tv_sec);

result = select(32, (fd_set *)&recv_fds, NULL, NULL, &timeout);
printf("select( ) returns: %d\n", result);

```

- Example Program: (Demo above code)

Example

- The **select()** test Prog file: [click here](#)
- Use this **sender** to send a **message** to the **select** program: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory
- To compile the **select1** test prog: **gcc -o select1 select1.c**
- To compile the **udp4-s3** test prog: **gcc -o udp4-s3 udp4-s3.c -lns1**
- To run: **select1 8000** on machine X
- To send a msg to **select** prog: **udp4-s3 X 8000**

Selecting (clearing/setting) the (socket) descriptors in general

- The *data type* of the descriptor sets `recv_fds`

- What is a Descriptor set:

- The variable `recv_fds` in the `select()` call:

```
select ( 32, &recv_fds, NULL, NULL, &timeout );
```

is the **descriptor set** (that is being *checked* for **readiness**)

- The **data type** of the `recv_fds` variable is:

```
fd_set recv_fds ;
```

The `fd_set` variable is:

- a **bit array** of size `FD_SETSIZE` (=1024) number of bits

- Each bit in the **bit array** variable represents **one entry** in the **descriptor table** of the process (= running program)

(Use: "man select" to find out more)

- In my **previous example**, I used a **short** bit array:

```
int recv_fds; // int is a bit array of 32 bits
```

This will **work** as long as you use **≤ 32 descriptors** (i.e., use less than around **30 sockets !!!**)

• Using `select()` with a `fd_set` variable

- How the `select()` function with a `fd_set` variable:

```
fd_set recv_fds; // bit array variable

int s = socket(PF_INET, SOCK_DGRAM, 0); // Socket that we want to
                                         // test for receive readiness

(1) clear all bits in the bit array variable recv_fds
(2) set the bit position s in the bit array variable recv_fds
(3) call:

select ( FD_SETSIZE, &recv_fds, NULL, NULL, timeout )
```

- Compare to my **simplified** usage (with an `int` variable):

```
int recv_fds; // Shorter bit array variable

int s = socket(PF_INET, SOCK_DGRAM, 0); // Socket that we want to
                                         // test for receive readiness

(1) clear all bits in recv_fds
    recv_fds = 0

(2) set the bit position s in recv_fds
    recv_fds = recv_fds | ( 1 << s )

(3) Call:
    select ( 32, &recv_fds, NULL, NULL, timeout );

(I need to use 32 to tell select( ) that it only need to
scan the first 32 entries in the descriptor table)
```

• Clearing, Setting and Testing *bits* in a bit array

- Bit array manipulation **macros** (defined inside `select.h`):

- `void FD_ZERO(fd_set *fdset)`

- Clear all bits in the "fdset" bit array variable

Example:

```
fd_set recv_fds;
FD_ZERO ( & recv_fds );

Result: recv_fds = 0000....000
```

- **void FD_SET(int s, fd_set *fdset)**

■ Set bit "s" in the "fdset" bit array variable to 1 (i.e., select "s" for checking)

Example:

```
fd_set recv_fs;
FD_ZERO ( & recv_fs );
FD_SET ( 2, & recv_fs );

Result: recv_fds = 0000....0100 // Bit position 2 is set
```

- **void FD_CLR(int fd, fd_set *fdset)**

■ Clear (reset) bit at position "fd" in the "fdset" bit array

- **int FD_ISSET(int fd, fd_set *fdset)**

■ Test the bit at position "fd" in the "fdset" bit array

Returns:

■ true (non-zero) if bit is set
■ false (= 0) if bit is not set

- The *general* usage of `select()` to check receive readiness of a socket

- Here is the *same demo* program re-written using a `fd_set` typed variable:

I have highlighted the changes

```
// I need to create a socket and bind it first.....

/* ---
   Create a socket
   --- */
s = socket(PF_INET, SOCK_DGRAM, 0);

/* ---
   Set up socket end-point info for binding
   --- */
in_addr.sin_family = AF_INET;           /* Protocol family */
in_addr.sin_addr.s_addr = htonl(INADDR_ANY); /* Wildcard IP address */
in_addr.sin_port = htons(atoi(argv[1])); /* Use any UDP port */

/* ---
   Bind socket s
   --- */
bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr))

/* =====
   How to use select
   ===== */
fd_set recv_fds;           /* I'm using a fd_set var now */
struct timeval timeout;    /* Time value for time out */

/* =====
   Select the bit for socket s
   ===== */
FD_ZERO(&recv_fds);        /* Clear all bits in recv_fds */
FD_SET(s, &recv_fds);      /* Set bit for socket s

/* =====
   Set timeout value to 5 sec
   ===== */
timeout.tv_sec = 5;         /* Timeout = 5 sec + 0 micro sec */
timeout.tv_usec = 0;        /* number of microseconds

int result;
```

```
printf("Now send a UDP message to port %s before %d sec !!!....\n",
      argv[1], timeout.tv_sec );

result = select(FD_SETSIZE, &recv_fds, NULL, NULL, &timeout);
printf("select( ) returns: %d\n", result);
```

- **Example Program:** (Demo above code)

Example

- The **modified select()** test Prog file: [click here](#)
- Use this **sender** to send a **message** to the **select** program: [click here](#)

How to run the program:

- **Right click** on link(s) and **save** in a scratch directory
- To compile the **select2** test prog: **gcc -o select2 select2.c**
- To compile the **udp4-s3** test prog: **gcc -o udp4-s3 udp4-s3.c -lssl**
- To run: **select2 8000** on machine X
- To send a msg to **select** prog: **udp4-s3 X 8000**

Timeout Programming using the `poll()` system call

- Timeout programming using the `poll()` system call

- The `poll()` system call:

- The `poll()` performs a **similar task** to `select()`:
 - `poll()` waits for **one** of a **set of sockets** to become **ready** to perform **read (receive) / write (send)** operation

- Differences:

- **How to** specify the **socket** to test for **readiness**
 - _____
 - **How to** to specify the **time out value**

- Header file for using `poll()`

- Header file for `poll()`:

```
#include <poll.h>
```

- The `poll()` function

- The `poll()` function **syntax**:

- Syntax:

```
#include <poll.h>          /* Header file */  
int poll( struct pollfd fds[],  nfds_t nfds,  int n_msec );
```

- Meaning of the parameters:**

```
fds      = an array of "struct pollfd"  
        This array contains the file descriptors (sockets) that  
        you want to check.  
  
(The pollfd structure will be explained later)  
  
n_fds   = the length of the "struct pollfd" array  
        (= number of descriptors to poll).  
  
n_msec  = the timeout value in milli-seconds  
  
Special timeout values:  
  
n_msec = 0 ==> poll() returns immediately with the poll results  
n_msec = -1 ==> poll() blocks until a requested event occurs
```

- Effect:

- The `poll()` function examines the **array of file/socket descriptors `fds[]`** for **`n_msec`** milli-seconds.

- If **some** descriptor (e.g., a socket) becomes **ready** before the **before** the **time out** value of **`n_msec`** time:

- `poll()` will return **immediately**, and:

- The **return value** is equal to the **number of descriptors (sockets)** that are **ready**
 - The **`fds[]`** array will contain **information** on **which** of the **descriptors/sockets** are **ready**

- If **after waiting** for **`n_msec`**, there are **not any descriptor (= socket) ready**:

- The `poll()` function will return
 - The **return value** will be equal to **0 (ZERO)**

- The **pollfd** (poll file descriptor) data structure

- The **struct pollfd** data structure is as follows:

```
struct pollfd
{
    int fd;          // A descriptor (e.g., socket)

    short events;   // This is a bit array
                    // Each bit represents a certain operation
                    // Set that bit to check for readiness
                    // E.g., one or the bit represent receive

    short revents;  // ready events
                    // This is also a bit array
                    // Each bit represents whether some operation is ready
}
```

Meaning of the parameters:

fd	= the (one) file descriptor/socket to be checked
events	= input bit array of events that will be checked
revents	= output bit array of events that are ready

The **types of events** are defined by the following **bit masks**:

POLLIN (= 00001) = ready for input (receive)	
POLLOUT (= 00010) = ready for output (send)	
Other less used functions:	
POLLRDBAND (= 00100) = high priority read data ready	
POLLWRBAND (= 01000) = high priority output (send) ready	
POLLERR (= 10000) = error detected	

- Examples on how the **data field** is used:

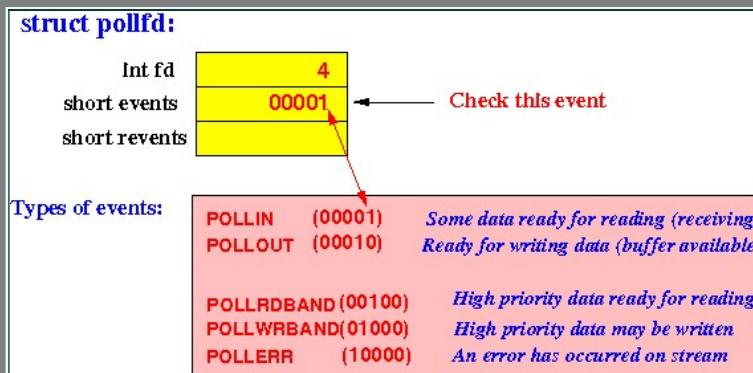
- Suppose you want to **check** if:

▪ the socket s has data available for reading (= receiving)
--

then you **set** the **values** as follows:

fds[i].fd = s;
fds[i].event = POLLIN;

Result:



- Checking for ready events in **poll()**

- Checking for readiness:

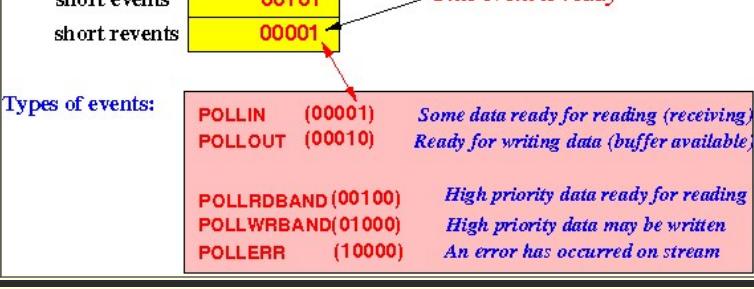
▪ The revents variable is an output field that contains the events that are ready

- Example:

- If the **file/socket #4** has some **normal priority data ready** (= available), then the **revents** variable will be **set** to:

struct pollfd:	<table border="1"> <tr> <td>Int fd</td> <td>4</td> </tr> <tr> <td>short events</td> <td>00101</td> </tr> </table>	Int fd	4	short events	00101
Int fd	4				
short events	00101				

This event is ready



- Setting the descriptor sets in the `pollfd` structure with sockets

- Example 1:

- We have **1 socket**:

```
int s = socket( PF_INET, SOCK_DGRAM, 0 );
```

in our program

- We want to **check** the **socket s** for **receive** readiness

Solution:

```
Define an pollfd arary of 1 element:

struct pollfd  fds[1];      // Array of 1 element
(The array element has index 0: fds[0])
```

```
Initialize the poll structure fd[0] to test s for read readiness:

fds[0].fd = s;                      // set descriptor fd to s
fds[0].events = 0;                  // Clear the bit array
fds[0].events = fds[0].events | POLLIN; // Set the POLLIN bit
```

```
Call poll( ) with timeout:

timeout = 100; // 100 msec = 0.1 sec

if ( poll ( fds, 1, timeout ) == 0 )
{
    // Time out
}
else
{
    // Ready
}
```

- Example 2:

- We have **2 sockets**:

```
int s1 = socket( PF_INET, SOCK_DGRAM, 0 );
int s2 = socket( PF_INET, SOCK_DGRAM, 0 );
```

in our program

- We want to **check** when **either socket s1 or s2** has **some message** that can be **received**

Solution:

```
Define a pollfd arary of 2 elements (because we have 2 sockets):

struct pollfd  fds[2];      // Array of 2 element
(The array elements are: fds[0] and fds[1])
```

```
Initialize the poll structure fd[0] to test s1 for read readiness:
```

```
    fds[0].fd = s1;                                // set descriptor fd to s1
    fds[0].events = 0;                             // Clear the bit array
    fds[0].events = fds[0].events | POLLIN; // Set the POLLIN bit
```

```
Initialize the poll structure fd[1] to test s2 for read readiness:
```

```
    fds[1].fd = s2;                                // set descriptor fd to s2
    fds[1].events = 0;                             // Clear the bit array
    fds[1].events = fds[0].events | POLLIN; // Set the POLLIN bit
```

```
Call poll( ):
```

```
    timeout = 100; // 100 msec = 0.1 sec
    if ( poll ( fds, 2, timeout ) == 0 )
    {
        // Time out
    }
    else
    {
        for ( i = 0; i < nfds; i++ )
        {
            if ( (fds[i].revents | POLLIN) == POLLIN )
            {
                // fds[i].fd is ready
            }
        }
    }
```

- Example: Timeout programming with `poll()` system call

- Here is the `select` example (discussed here: [click here](#)) re-written using `poll()`:

```
// I need to create a socket and bind it first....
/* ---
   Create a socket
   --- */
s = socket(PF_INET, SOCK_DGRAM, 0);

/* ---
   Set up socket end-point info for binding
   --- */
in_addr.sin_family = AF_INET;           /* Protocol family */
in_addr.sin_addr.s_addr = htonl(INADDR_ANY); /* Wildcard IP address */
in_addr.sin_port = htons(atoi(argv[1])); /* Use any UDP port */

/* ---
   Bind socket s
   --- */
bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr))

/* =====
   How to use poll( )
   ===== */
struct pollfd  fds[1];                  // 1 socket
int      timeout = 5000;                // 5000 msec = 5 sec

/* =====
   Set up fds[0] to monitor socket s
   ===== */
fds[0].fd = s;                         // Test socket s
fds[0].events = 0;                     // Clear events
fds[0].events = fds[0].events | POLLIN; // Set INPUT bit

int result ;

printf("Now use udp4-s3 and send a UDP message to port %s before %d sec\n",
       argv[1], timeout/1000 );

/* =====
   Here we go...
   ===== */
result = poll( fds, 1, timeout );
printf("poll( ) returns: %d\n", result);
```

- The **poll()** test Prog file: [click here](#)
- Use this **sender** to send a **message** to the **select** program: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory
- To compile the **poll1** test prog: `gcc -o poll1 poll1.c`
- To compile the **udp4-s3** test prog: `gcc -o udp4-s3 udp4-s3.c -lns1`
- To run: `poll1 8000` on machine X
- To send a msg to **select** prog: `udp4-s3 X 8000`

Intro to TCP communication

- Properties of TCP communication

- Properties of TCP communication:

- Reliable:
 - Data sent using a **TCP socket** is received **once** and in the transmission order]
 - Stream oriented
 - Data transmitted using TCP can be received in **different chunk sizes** than transmitted !!!

Graphically:

See: [click here](#)

- Note for CS450 students:

- A **TCP socket** behaves **exactly the same** as a **UNIX pipe**!

Except that:

- You use **one (= same) descriptor** for reading/writing !!!

I.e.:

- A **TCP socket** is **bi-directional**

- TCP and Network programs

- Comment:

- TCP sockets are the **most commonly used communication end point** in **network programs**

because:

- Retransmissions (to provide **reliability**) is handled by the **TCP protocol** itself.

(So the **programmer** can **focus** on **programming** the **application** itself....)

- Virtually **every** network program is written using **TCP sockets**

(**UDP sockets** are only used by **network research projects** that experiment with **new network transmission algorithms**)

TCP sockets and TCP connections

- Creating a TCP socket

- Create a **TCP socket**:

```
int s;           // Descriptor index.  
  
s = socket(AF_INET, SOCK_STREAM, 0);  
  
          // Creates a TCP socket in the per-process descriptor table  
// Returns the index in the per-process descriptor table
```

- TCP sockets and TCP connections

- **TCP socket**:

- **TCP socket** = a **end point** of **communication using the TCP protocol**
 - Unlike UDP, the **TCP protocol** does **not only** use the **TCP socket** to **identify** a **programming communication endpoint**

- **TCP connection**:

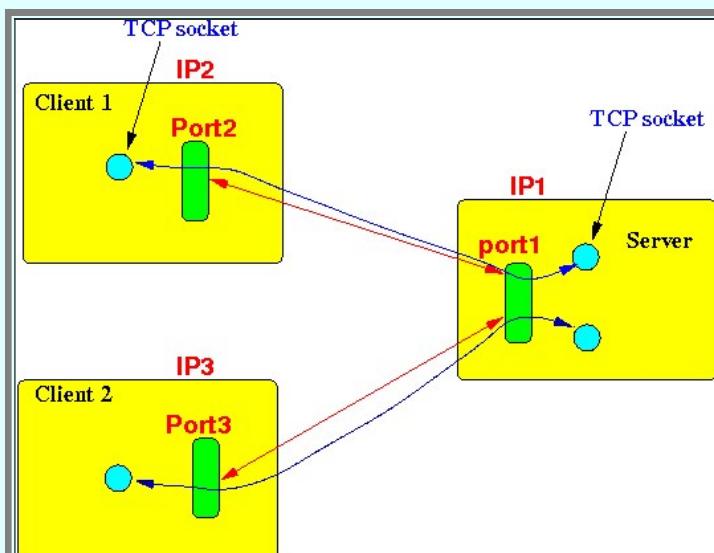
- **TCP connection** = a **communication link** between **2 network programs** that uses **TCP sockets**
 - Network programs communicate with one another using **TCP connections**
 - A **TCP connection** is **uniquely identified** by:
 - (srcIP, srcPort#, destIP, destPort#)

- Multiple TCP connections using the **same** (IP-address, Port#)

- Fact:

- **Because:**
 - A **TCP connection** (= **communication pathway** between **2 network programs**) is **identified** by:
 - (srcIP, srcPort#, destIP, destPort#)
 - it is **possible** to **form** the following **different TCP connections**:
 - TCP connection 1: (IP1, port1, IP2, port2)
 - TCP connection 2: (IP1, port1, IP3, port3)
 - using the **same** (IP1, port1)

Graphically:



This is the **situation** where a **web-server** running on the **host IP1** is **servicing 2 web clients**:

- One **client** running on **host IP2**
- One **client** running on **host IP3**

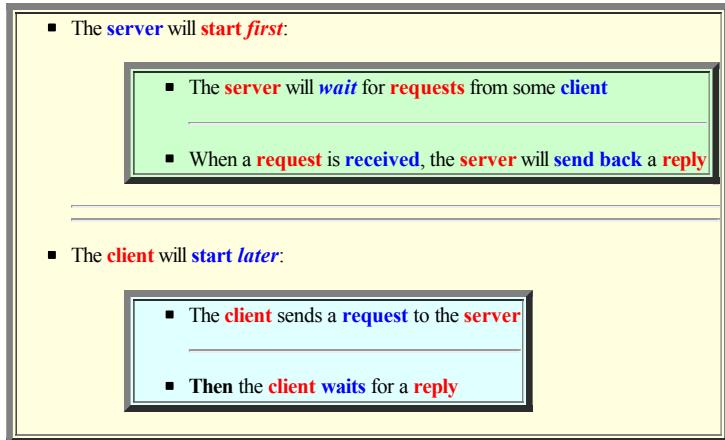
In other words:

- **One TCP port** in the **Operating System kernel** can support ***multiple* TCP connections** for the ***same IP address***

Programming with TCP

- **TCP programming model**

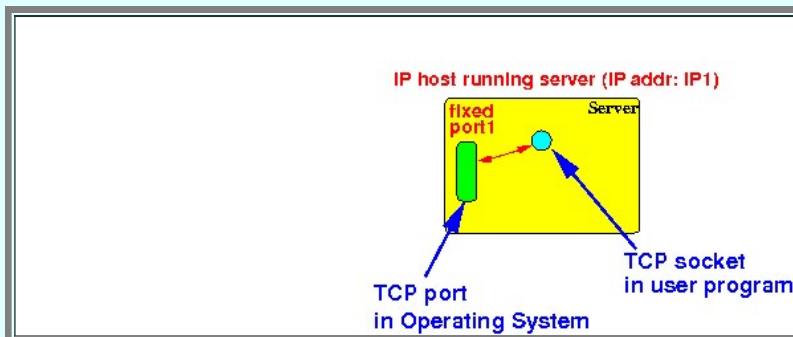
- Programming using TCP will *always* use the **client/server paradigm**:



- **How to use TCP in client/server programming**

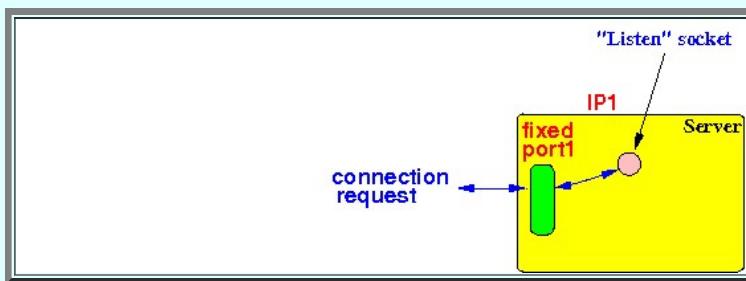
- Steps in writing a **client/server network program** using TCP:

1. The **server program** first creates a **TCP socket** and **binds** it to a **fixed TCP port#**:



(We have learned this step before in **UDP programming**)

2. Then the **server program change** the **socket** to **listen socket** that is used to **receive connection requests**:



The **listen socket** is a **special type of socket** that **execute** the **"3-way handshake"** algorithm to **establish** a **reliable connection**

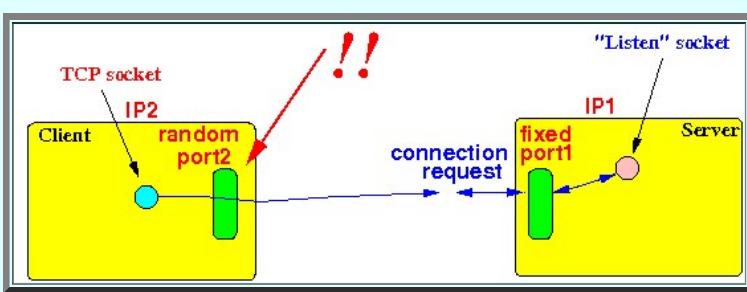
- The **server** will now **wait** for "**client connection requests**"

(This is done using the `accept(...)` system call - discussed later)

3. A **client program** first **create** a **TCP socket**:



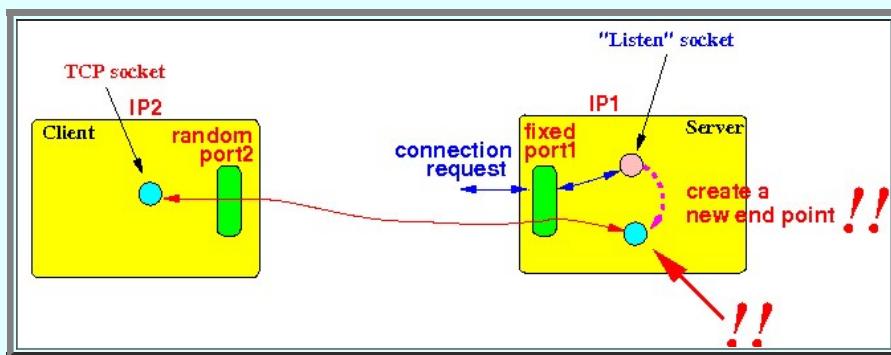
Bind the **socket** to an **arbitrary TCP port #** and use the **socket** to make a **connection request** to the **server program** (at the **listen socket**):



The **client program** use the `connect()` system call to do this

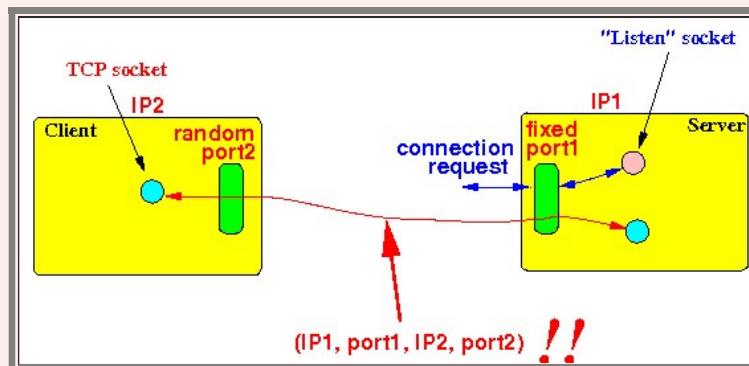
(The **TCP connection process** will establish a **pair** of send/ACK sequence numbers for the **sliding window algorithm**)

- When the **processing** of the **connection request** is **complete**, the **TCP module** at the **server program** will **create a new TCP socket** for the **TCP connection**:

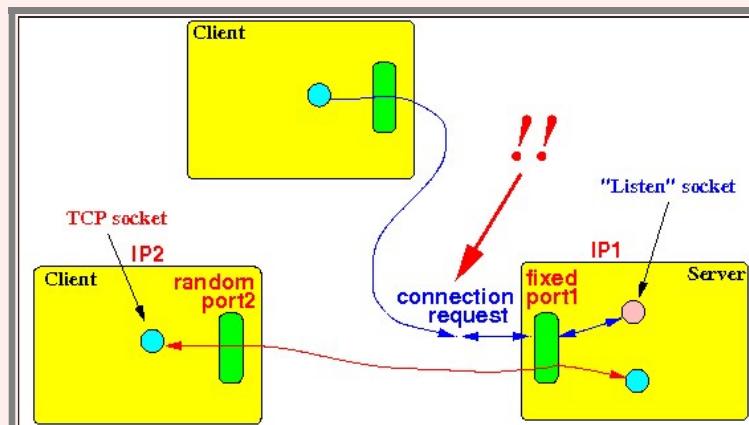


5. Result:

- The **client program** and the **server program** can **communicate** with each other through the **TCP connection (IP1, port1, IP2, port2)**:



- And*, the **server** can **still** accept **other connection requests** using the **listen TCP socket**:



The **listen socket** remains **free** to receive **new connection requests** from **other clients** !!!

- This is **how** a **web server** or **SSH server**, and so on... operates

Programming the server using TCP

- Programming a TCP server application

- Recipe (steps) for coding a TCP server application:

- Create a TCP socket

```
int s = socket( AT_INET, SOCK_STREAM, 0 );
```

This **socket** is *only* used for:

- Receiving TCP connection requests from clients !!!

- Bind the **socket s** to a *fixed* TCP port#:

```
struct sockaddr_in in_addr;  
  
in_addr.sin_family      = AF_INET;  
in_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
in_addr.sin_port        = htons(SrvPortNumber); // a fixed port#  
  
bind(s, (struct sockaddr *)&in_addr, sizeof(in_addr))
```

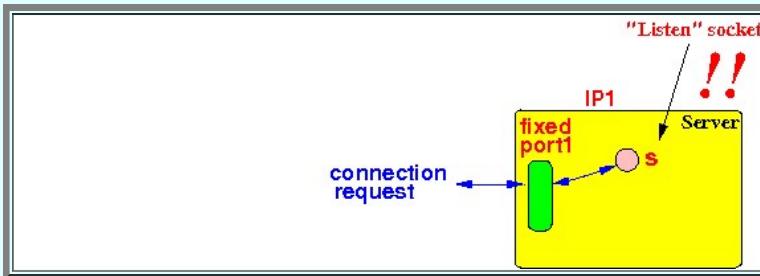
Result:

- TCP messages can be received on this Internet network address

- Enable the **socket s** to receive client's connection requests:

```
listen( s, #pending-connect-requests );
```

Result:

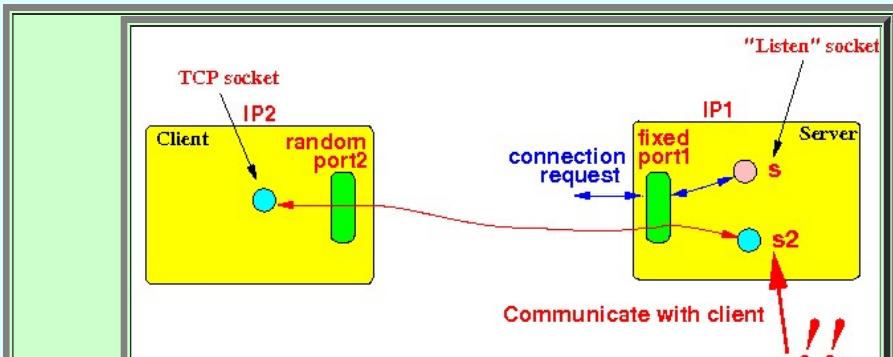


More details on the `listen()` function in *later*

- Establish TCP connections with clients:

```
int s2; // Another socket variable  
  
s2 = accept( s, &in_addr, &length );
```

Result:



- The socket **s2** is used to communicate with a **client**

- The socket **s** will *still* be used to establish TCP connections with (other) clients

- The **listen** system call

- Syntax of the **listen()** function:

```
int listen(int s, int backlogLength)

s           = a socket descriptor
backlogLength = max. number of pending requests allowed

listen returns

0 for success,
-1 for error
```

- Effect of the **listen(s, ...)** call:

- **listen(s, ...)** enables the **socket s** to **accept** incoming **connection requests** sent by **TCP client applications**

Note:

- A **TCP socket s** does *not accept* client's **connection requests**
 -
 - You must **enable** the **connection processing** with the **listen(s, ...)** call

- The **accept** system call

- Syntax of the **accept()** function:

```
int accept ( int s, struct sockaddr *cli_addr, socklen_t *addrlen)

Input:
      s           = a listen socket

Output:
      cli_addr    = network address (will contain (IPAddr, Port#) of client
      addrlen     = length of addr

Return value:
      a new TCP socket descriptor

Note: use the returned TCP socket to
      send messages to client
      receive messages from client
```

- Effect of **accept()** function:

- The **accept()** function will **block** until a **TCP connection** is **successfully established**:

- an **incoming connection request** is **received**
 -
 - a **consistent TCP state** has been **negotiated** (**send/ACK numbers**)

- The **TCP connection establishment** will **perform** the **following** in the **OS kernel**:

- Create a **new descriptor entry** in the **per-process descriptor table** for a **TCP socket** containing:
 - The **state (send/ACK numbers)** of the **TCP connection**
 -
 - The **send/receive buffer spaces** for the **new TCP connection**

- The **accept()** function will **return** the **(socket) descriptor**

- The **server** must use the **return descriptor** to **send/receive data** to/from the **client**
(Because this **socket** contains the **TCP state** for **reliable communication**)

- Structure of a **server TCP application** written in **C program code**:

```

int s_listen;           // The listen socket
int s_data;             // The data socket

struct sockaddr_in in_addr;    // Help structures to form netw address
struct sockaddr_in client;
int length;

/* -----
   Create the TCP listen socket
----- */
s_listen = socket( AF_INET, SOCK_STREAM, 0 );      // Create a TCP socket

/* -----
   Bind the socket to a network address
   This makes the socket identifiable by (IPAddr, Port#)
----- */
in_addr.sin_family = AF_INET;
in_addr.sin_addr.s_addr = htonl(INADDR_ANY);
in_addr.sin_port = htons(SrvPortNumber);

if ( bind(s_listen, (struct sockaddr *)&in_addr,      // Bind socket to (IP-addr, Port#)
           sizeof(in_addr)) == 0 )
{
    perror("Bind: ");
    exit(1);
}

/* -----
   Enable socket s_listen to accept connect requests
----- */
listen(s_listen, 5);      // Change socket to a "listen" socket

/* -----
   Accept TCP connection calls
----- */
while ( 1 )
{
    length = sizeof(struct sockaddr_in);

    /* =====
       Server program waits for client's connection requests
===== */
    s_data = accept( s_listen, &client, &length);      // s_data = data socket

    // The accept( ) function will a
    // data socket of the TCP connection
    // establish with the client program

    /* =====
       Now you can send and receive data
       using the s_data socket
===== */

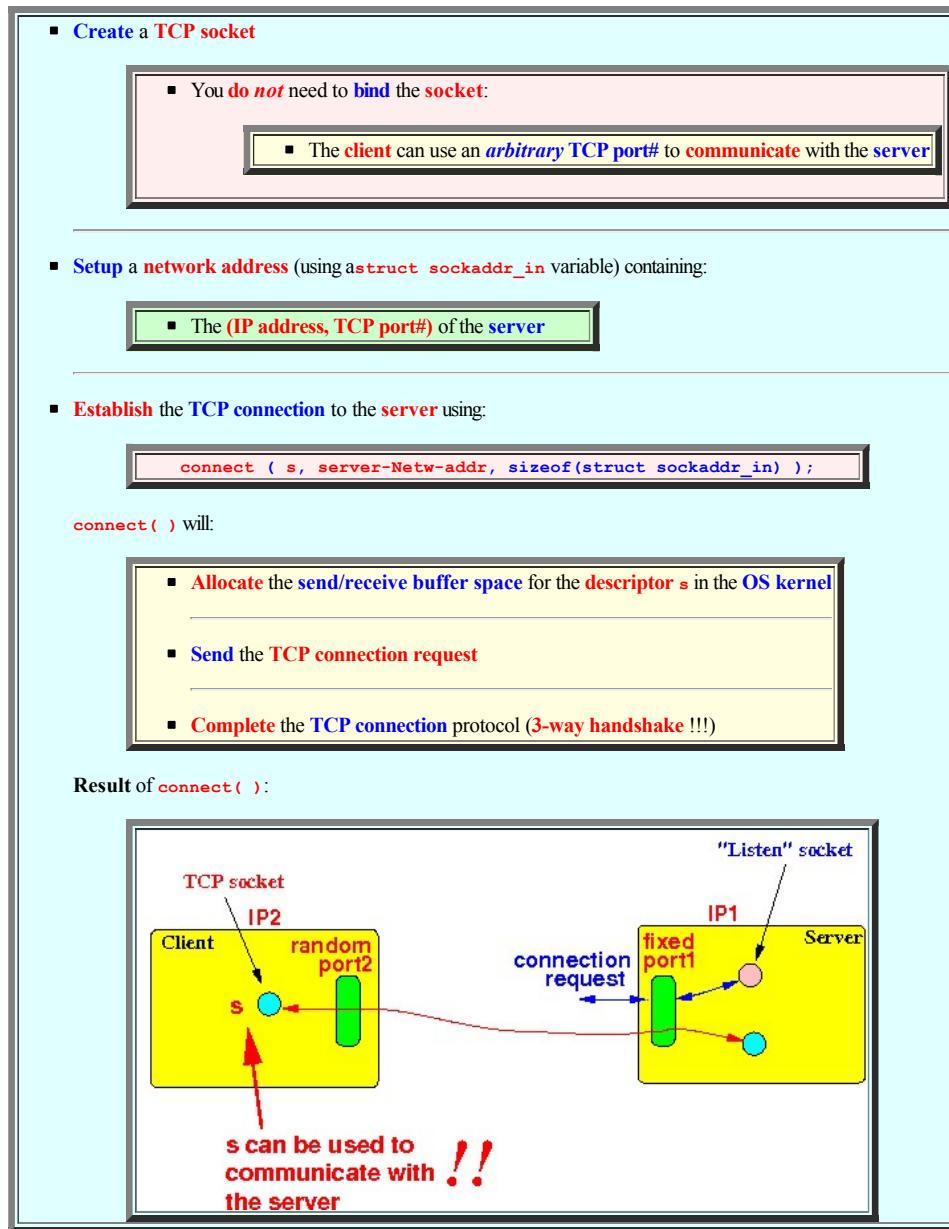
    read ( s_data, ....)    to receive client's requests
    write( s_data, ....)    to send data to client
    close( s_data)          to disconnect
}

```

Programming the client using TCP

- Structure of a TCP client application

- Recipe (steps) for coding a TCP *client* application ("client-side" programming):



- The `connect()` function

- Syntax of the `connect()` system call:

```
int connect(int s, struct sockaddr *server, length)

Input:
    s      = an unbounded socket !!!
    server = a "struct sockaddr_in" structure containing
             the IP-address and the port# of the server
    length = must be equal to "sizeof(struct sockaddr_in)"

Return value:
    0      = success
    -1     = failed
```

- Effect:

- The `connect()` system call will establish a *reliable connection* between the **client** and the **server application** given at the **Internet address** in `server`
 - The `server` structure must contains the *listen port* of the **server program**
- The `connect()` function will **reserve** an *unused TCP port* in the **Operating System** and **bind** this TCP port to the **descriptor entry s**

◦ **Result:**

- After the **connect()** call is **completed**, the **client application** can use the **socket variable s** to:
 - **send** messages (= requests) to the **server's data TCP port**
 - **receive** messages (= replies) from the **server's data TCP port**

• Structure of a TCP client application program

◦ **Structure** of a **TCP client application**:

```

int s;                                // Socket variable
struct sockaddr_in server;             // Network addr for server

/* -----
   Create TCP socket
   ----- */
s = socket( AF_INET, SOCK_STREAM, 0 );    / Create a TCP socket

/* -----
   Setup the server's Network Address
   ----- */
server.sin_family      = AF_INET;
server.sin_addr.s_addr = htonl(SrvIP-address);
server.sin_port         = htons(SrvPortNumber);

/* -----
   Make a connection request to the server
   (The server must be in an accept() call to be successful)
   ----- */
if ( connect(s, server, sizeof(struct sockaddr_in)) == 0 )
{
  perror("connect: ");
  exit(1);
}

// After this, the client can use the TCP socket s to communicate
// (to the data TCP port) with server
//
// Use:

  read(s, ....)      to get server's replies
  write(s, ....)     to send data to server
  close(s)           to disconnect

```

Sending/receiving data with TCP

- Sending messages on TCP sockets

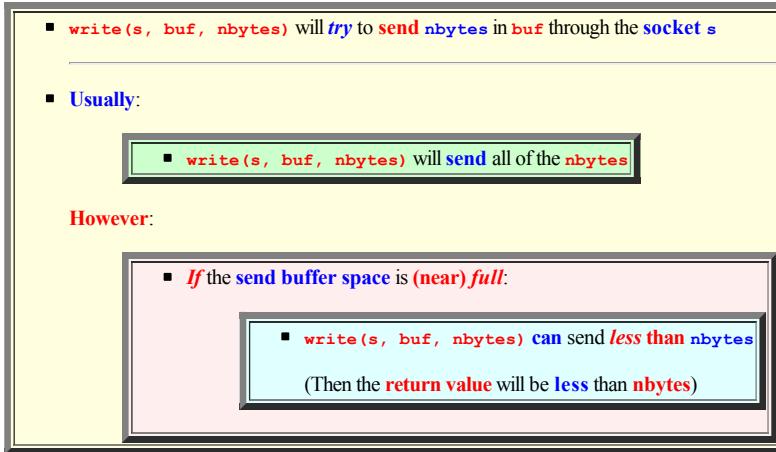
 - Sending (= writing) with a TCP socket s:

```
int write(int s, void *buf, int nbytes)

Input parameters:
    s      = the TCP socket
    buf    = buffer address containing the data to be sent
    nbytes = number of bytes in buffer "buf" to send

Return value:
    -1    = error
    >=0   = the actual number of bytes sent
```

Effect:



- A commonly used TCPsend() function

 - The following function is guaranteed to send all nBytes:

```
int TCPsend( int s, char *buf, int nBytes )
{
    int k, n;

    k = 0;           // No bytes sent yet.....

    /* -----
       Call write( ) as long as we did not send nBytes.....
    ----- */
    while ( k < nBytes )
    {
        n = write( s, &buf[k], nBytes-k ); // Send "remaining" nBytes-k bytes

        /* =====
           Check if there was an error....
        ===== */
        if ( n == -1 )
            return -1;                  // Return error.....

        k = k + n;                   // We sent n more bytes, add to k
    }

    return nBytes;           // Success...
}
```

- Receiving messages on TCP sockets

 - Receiving (= reading) data with a TCP socket s:

```
int read(int s, void *buf, int nbytes)

Input:
    s      = TCP socket
    buf    = address of the buffer to receive read data
    nbytes = # bytes that we wish to receive

Return value:
    -1    = error
    >=0   = the actual number of bytes transferred
```

Effect:

- `read(s, buf, nbytes)` will **block** until **some data** is **available**
- When **data** is **available**, `read(s, buf, nbytes)` will **try** to:
 - Transfer **nbytes** from the **receive buffer** of **socket s**
 - **Store the data** in the **buf** variable
- **Usually:**
 - `read(s, buf, nbytes)` will **transfer exactly nbytes**
- **However:**
 - When the receive buffer is **near empty**, the `read` operation **may return** with **less than nbytes**.

• A commonly used TCPRecv() function

- The following **function** is **guaranteed** to **receive** exactly **nBytes**:

```
int TCPRecv( int s, char *buf, int nBytes )
{
    int k;

    k = 0;                      // No bytes received yet.....

    /* -----
     * Call read( ) as long as we did not receive nBytes.....
     */
    while ( k < nBytes )
    {
        n = read( s, buf[k], nBytes-k ); // Receive "remaining" nBytes-k bytes

        /* =====
         * Check if there was an error...
         */
        if ( n == -1 )
            return -1;                  // Return error.....
        else
            k = k + n;                // We sent n more bytes, add to k
    }

    return nBytes;
}
```

• Closing a TCP connection

- **Closing a TCP connection:**

```
int s;      // TCP socket used to communicate
close(s);
```

• For CS450 students

- **Programming remark for CS450 students:**

- A **TCP socket** behaves exactly like a **pipe** or an **file descriptor** of an open file !
 - The `read()/write()` call can **return** with **less than** the **specified amount of data** when we reach the **end of the file !!!**

The **similarity** is **obvious** in the **behavior**:

- a **pipe** or an **opened file** in **UNIX** is a **stream-oriented device**:
 - the **unit (smallest amount) of data** that can be **read/written** is **one byte**

Programming a *serial* server (easy to understand, no self-respecting programmer will do it)

- *Serial server: serving one client at a time*

- ### ◦ **Serial server:**

- A **serial server** can **attend to one client** at a time
 - While **processing** the client, the **TCP server** **can not accept** a **new connection request**

(Because the **server** is **not able** to execute the `accept()` call....)

- **Structure** of a *serial TCP server*:

- Example 1: "print to terminal" server

- This processing loop will **print** each character that is **received**

```
/* -----
   Processing loop
----- */  
while (read(s2, &c, 1) > 0) // Receive 1 character  
{  
    putchar(c);           // Print it  
    fflush(stdout);  
}
```

```
close(s2);
```

- The entire **TCP server** program:

```
main(int argc, char *argv[])
{
    int s1; /* s1 = listen socket */
    int s2; /* s2 = connect socket */
    struct sockaddr_in me; /* My own Internet address */
    struct sockaddr_in myname; /* Used to retrieve my TCP port */
    int myname_len;

    struct sockaddr_in from; /* Internet address from client */
    int len; /* length associated with 'from' */

    char c; /* Used to get data from connection */

    if (argc == 1)
        { printf("Usage: %s port\n", argv[0]);
        exit(1);
    }

    /* -----
    Create a TCP socket
    ----- */
    s1 = socket(AF_INET, SOCK_STREAM, 0);

    /* -----
    Set up Internet address (IPaddr, Port#) to bind the socket
    ----- */
    me.sin_family = AF_INET; /* Put in the family */
    me.sin_port = htons(atoi(argv[1])); /* Put in the port number */
    me.sin_addr.s_addr = htonl(INADDR_ANY); /* Use wildcard address */

    /* -----
     Make TCP socket "identifiable"
    ----- */
    if (bind(s1, (struct sockaddr *) &me, sizeof(me)) < 0)
    { perror("bind:");
    exit(1);
    }

    /* =====
    Change s1 into a "listen" socket
    (with max 5 pending connection reqs)
    ===== */
    listen(s1, 5);

    /* -----
    Keep the server running forever....
    ----- */
    while (1)
    {
        len = sizeof(from);

        /* =====
        Make the server program wait for a connection
        from a client...
        ===== */
        if ( (s2 = accept(s1, (struct sockaddr *) &from, &len)) < 0)
        {
            perror("accept: ");
            exit(1);
        }

        /* =====
        Use socket s2 to get data from client
        and print it
        ===== */
        while (read(s2, &c, 1) > 0)
        {
            putchar(c); // Print each character read
            fflush(stdout);
        }

        close(s2);
        printf("\n\n...Done...\n\n");
    }
}
```

- The client simply sends a **serie of characters** to the **server**:

```
/* =====
The socket creation and connection to server
have been omitted for brevity
===== */

char c;

while ((c = getchar()) != EOF) // Read a character and store in c
{
    write(s, &c, 1); // Send the character to server
}
```

- **Example Program:** (Demo above code)

Example

- The **server** Prog file: [click here](#)

- The **client** Prog file: [click here](#)

How to run the program:

- Right click on link and **save** in a scratch directory
- To compile: `gcc -o tcp1-cli tcp1-cli.c`
`gcc -o tcp1-serv tcp1-serv.c`
- To run: on machine **X**: `tcp1-serv 8000`
on machine **Y**: `tcp1-cli X 8000`

• Example 2: addition server

- The **server** will **read 2 integers** and then **send the sum** back to the **client**:

```
/* =====
   This is the processing loop of the server
===== */
int x, y, z;

while ( 1 )
{
    int xl, yl, zl;

    if ( read( s2, &xl, 4) <= 0 ) // An int is 4 bytes
        break;

    x = ntohl(xl);           // Convert to host order

    if ( read( s2, &yl, 4) <= 0 )
        break;

    y = ntohl(yl);           // Convert to host order

    z = x + y;
    printf("Client data: %d, %d, send sum = %d\n", x, y, z);

    zl = htonl(z);
    write ( s2, &zl, 4 );
}

close(s2);
```

- The **client** will send the **server 2 integer values** and then **receive back an integer (sum)** and prints it:

```
/* =====
   The socket creation and connection to server
   have been omitted for brevity
===== */

int x, y, z;

while ( 1 )
{
    printf("Enter first number (0 to exit): ");
    scanf("%d", &x);           // Read first integer

    if ( x == 0 )
    {
        close(s);
        exit(0);
    }

    printf("Enter second number: ");
    scanf("%d", &y);           // Read second integer

    int xl, yl, zl;

    xl = htonl(x);           // Convert them to network byte order first
    yl = htonl(y);

    write(s, &xl, 4);         // Send the integers
    write(s, &yl, 4);

    /* =====
       Receive the answer from the server
    ===== */
    read(s, &zl, 4);
    z = htonl(zl);           // Convert answer to host byte order

    printf("%d + %d = %d\n", x, y, z); // Print it
}
```

Example

- **Example Program:** (Demo above code)

- The **server** Prog file: [click here](#)
- The **client** Prog file: [click here](#)

How to run the program:

- Right click on link and **save** in a scratch directory
- To compile: `gcc -o tcpla-cli tcpla-cli.c`
`gcc -o tcpla-serv tcpla-serv.c`
- To run: on machine **X**: `tcpla-serv 8000`
on machine **Y**: `tcpla-cli X 8000`

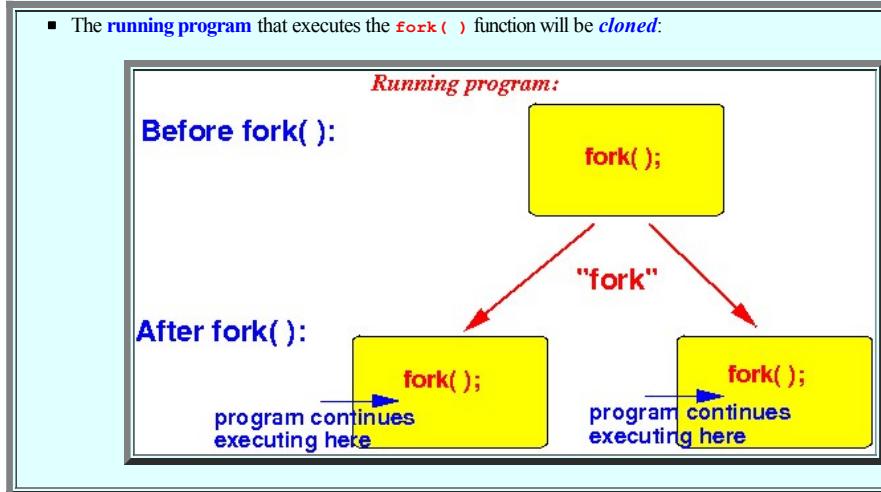
The fork() system call

- The `fork()` system call (from CS450)

- Syntax of the `fork()` system call:

```
int fork()
```

Effect:



- Example:

```
int main(int argc, char *argv[])
{
    printf("Hello World !\n");
    fork();
    printf("Good-bye World !\n");
}
```

Output:

```
Hello World !
Good-bye World !           // Printed twice !!!
Good-bye World !
```

Example

- Example Program: (Demo above code)

- Prog file: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory
 - To compile: `gcc fork1.c`
 - To run: `a.out`

- Using `fork()`

- Note:

- It may sound **stupid** at first to **clone** a **running program** because:
 - **why** on earth would you want to run the **same program twice** ???

- Using the `fork()` system call:

- The **is** a **way** (discuss soon) to make the **2 identical programs** continue **2 different ways**:
 - Use an **if-statement**
 - Then:

- One copy of the program will execute the `if`-part
- The other copy of the program will execute the `else`-part

- Parallel server

- Parallel server:

- Parallel server = a **server** than can **process multiple connection requests** from **clients** at the **same time**

- Structure of a **parallel server**:

- Make a **clone** of **yourself** (= the server program)
 - Have the **clone** provide **service** to the **client**
 - In the mean time, **you** (= the original server program) can **accept** other **connection requests** made by **other clients** !!!

- Parent and child process

- Parent process and child process:

- Parent process = the **original running program** that called the `fork()` function
 - Child process = the **clone** of the **original running program**

- Very important fact:

- The **return value** of the `fork()` function is **different**:
 - For the **parent process**, and
 - for the **child process**
 - The **return value** of `fork()` for the **parent process** is:
 - The **process id ($\neq 0$)** of the **child process**
 - The **return value** of `fork()` for the **child process** is:
 - **0 (always)**

- Example:

```
int main(int argc, char *argv[])
{
    printf("Hello World !\n");

    int i = fork();

    printf("i = %d\n", i);           // Show the return value of fork()
}
```

Sample output:

```
Hello World !
i = 10814      <-- This line was printed by the parent process
i = 0          <-- This line was printed by the child process
```

- Example Program: (Demo above code)

Example

- Prog file: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory

- To compile: `gcc fork2.c`
- To run: `a.out`

- How to make *parent* process and *child* process run *different* code

- `fork()` is *always* used inside an `if`-statement:

```
if ( fork( ) != 0 )
{
    code executed by the parent process
}
else
{
    code executed by the child process
}
```

- Example:

```
int main(int argc, char *argv[] )
{
    printf("Hello World !\n");

    if ( fork( ) != 0 )
    {
        printf("Hello, I am the parent process\n");
    }
    else
    {
        printf("Hello, I am the child process\n");
    }
}
```

Example

- Example Program: (Demo above code)

- Prog file: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory
- To compile: `gcc fork3.c`
- To run: `a.out`

- A more **advanced** example:

- The **parent process** will keep created **child processes** when you enter a **line**:

```
int main(int argc, char *argv[] )
{
    char line[1000];

    while ( 1 )
    {
        printf("Enter a line:\n");
        scanf( "%[^\\n]", line );
        getchar();

        if ( fork( ) != 0 )
        {
            printf("Parent process: I will go read a line again\n");
        }
        else
        {
            /* -----
            Code that is executed by the child process
            ----- */
            printf("Child process... your line is:    %s\n", line);
            printf("Child process... here is it again: %s\n", line);
            printf("Child process: Done !!\n");
            exit(0);                                // Child process exits !!
        }
    }
}
```

Example

- Example Program: (Demo above code)

- Prog file: [click here](#)

How to run the program:

- Right click on link(s) and **save** in a scratch directory
- To compile: `gcc fork4.c`

■ To run: **a.out**

Programming a *parallel* server with fork()

- *Parallel* server: serving multiple clients

- Parallel server:

- A **parallel server** can **attend to many client** at the same time
- Actually, the **reality** is that a **parallel server** make a **clone** of itself to **attend to one client** at a time...
- The **next request** can **accepted** by the **original**, while a **clones** service the **other requests**

- Coding a parallel server:

```
/* -----
   Create a TCP socket
----- */
s1 = socket(AF_INET, SOCK_STREAM, 0);

.....
/* -----
   Bind to a network address so that the server can listen
   for connection requests
----- */
if ( bind(s1, (struct sockaddr *) &server, sizeof(server) ) < 0 )
{ perror("bind:");
  exit(1);
}

/* -----
   Listen for connection requests on s1
----- */
listen(s1, 5);

while (1)
{
  len = sizeof(from);

  /* -----
     Accept connection requests
----- */
  if ( (s2 = accept(s1, (struct sockaddr *) &from, &len)) < 0 )
  { perror("accept: ");
    exit(1);
  }

  /* -----
     You arrive here when the server has accepted
   a connection request from one client

   Use s2 to communicate with the client
----- */

  if ( fork() == 0 )
  {
    /* -----
       Child process (the clone)
----- */
    close(s1);      // Don't need to listen

    /* -----
       Processing loop
----- */
    while ( read( s2, buf, .... ) > 0 )           | Change this
    {
      |          | to do
      process data in buf | other tasks
      write( s2, reply, ... ); | - - +
    }

    close(s2);
    exit( 0 );
  }
  else
  {
    /* -----
       Parent process (the original)
----- */
    close(s2);

    // We loop back up and accept the next connection request
  }
}

}
```

- Example 1: "null" server

- The **child process** will **print** each character that is **received**:

(While the **parent process** goes back and **executed** another **accept** call)

```
while (1)
{
  s2 = accept(s1, (struct sockaddr *) &from, &len));
```

```

/* -----
   fork()
----- */
if ( fork() == 0 )
{
    /* =====
       Code executed by the child process
       ===== */
    close(s1);           // Child does not use Listen socket

    /* =====
       Use data socket and comm. with client
       ===== */
    while ( read(s2, &c, 1) > 0 )
    {
        putchar(c);
        fflush(stdout);
    }

    close(s2);
    printf("\n\n...Done...\n\n");

    exit(0);           // Child process exits !!
}
else
{
    close(s2);
}
}

```

- **Example Program:** (Demo above code)

Example

- The **parallel server** Prog file: [click here](#)
- We can use the **same client** Prog file: [click here](#)

How to run the program:

- Right click on link and **save** in a scratch directory
- To compile: `gcc -o tcp1-cli tcp1-cli.c`
`gcc -o tcp2-serv tcp2-serv.c`
- To run: on machine X: `tcp2-serv 8000`
on machine Y: `tcp1-cli X 8000`
on machine Z: `tcp1-cli X 8000`

Both client program can send to the **server** !!!

- **Example 2: addition server**

- The **child process** will **read 2 integers** and then **send** the **sum** back to the **client** (until there is no more data):

(While the **parent process** goes back and **executed** another **accept** call)

```

while (1)
{
    s2 = accept(s1, (struct sockaddr *) &from, &len);

    if ( fork() == 0 )
    {
        /* =====
           Code executed by the child process
           ===== */
        close(s1);           // Child does not use the Listen socket

        /* =====
           Use data socket and comm. with client
           ===== */
        while ( 1 )
        {
            int x1, y1, z1;

            if ( read( s2, &x1, 4) <= 0 )
                break;

            x = ntohl(x1);           // Convert to host order

            if ( read( s2, &y1, 4) <= 0 )
                break;

            y = ntohl(y1);           // Convert to host order

            z = x + y;
            printf("Client data: %d, %d, send sum = %d\n", x, y, z);

            z1 = htonl(z);
            write ( s2, &z1, 4 );
        }

        close(s2);
        printf("\n\n...Done...\n\n");
        exit(0);
    }
    else
    {

```

```
        close(s2);  
    }  
}
```

- **Example Program:** (Demo above code)

Example

- The **parallel addition server** Prog file: [click here](#)
- We can use the **same client** Prog file: [click here](#)

How to run the program:

- **Right click** on link and **save** in a scratch directory
- To compile: `gcc -o tcp1a-cli tcp1a-cli.c`
`gcc -o tcp2a-serv tcp2a-serv.c`
- To run: on machine **X**: `tcp2a-serv 8000`
on machine **Y**: `tcp1a-cli X 8000`
on machine **Z**: `tcp1a-cli X 8000`

Both client program can send to the **server** !!!